

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
{
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE *f;
```

```
for(i=0; i<MAXPAROLA; i++)
    freq[i]=0;
```

```
if(argc != 2)
```

```
{
    fprintf(stderr, "ERRORE, serve un parametro con il nome del file\n");
    exit(1);
}
```

```
f = fopen(argv[1], "r");
if(f==NULL)
```

```
{
    fprintf(stderr, "ERRORE, impossibile aprire il file %s\n", argv[1]);
    exit(1);
}
```

```
while( fgets( riga, MAXRIGA, f ) != NULL )
```



# System and Device Programming

## UNIX Signals

Stefano Quer

Dipartimento di Automatica e Informatica

Politecnico di Torino

## Definition

### ❖ A **signal** is

- A software interrupt
- An asynchronous notification sent, by the kernel or by another process, to a process to notify it of an event that occurred

### ❖ Signals

- Allow notify asynchronous events
  - Such as the occurrence of particular events (e.g., error conditions, memory access violations, calculation errors, illegal instructions, etc.)
- Can be used as a limited form of inter-process communication

## Characteristics

- ❖ Available from the very first versions of UNIX
  - Originally managed in an unreliable way
- ❖ Standardized by the POSIX standard, they are now stable and relatively reliable
- ❖ Each signal has a name
  - Names start with **SIG...**
  - The file **signal.h** defines signal names
    - Unix FreeBSD, Mac OS X and Linux support 31 signals
    - Solaris supports 38 signals

The shell command **kill -l** displays a complete list of signals

## Main signals

Name	Description
SIGABRT	Process abort, generated by system call abort
SIGALRM	Alarm clock, generated by system call alarm
SIGFPE	Floating-Point exception
SIGILL	Illegal instruction
SIGKILL	Kill (non maskable)
SIGPIPE	Write on a pipe with no reader
SIGSEGV	Invalid memory segment access
SIGCHLD	Child process stopped or exited
<b>SIGUSR1 SIGUSR2</b>	<b>User-defined signal 1/2 default action = terminate the process Available for use in user applications</b>

# Signal management

❖ Signal management goes through three phases:  
signal generation, signal delivery, reaction to a signal

## ➤ Signal generation

- When the kernel or a source process causes an event that generate a signal

## ➤ Signal delivery

- A not yet delivered signal remains pending
- At signal delivery a process executes the actions related to that signal
- The lifetime of a signal is from its generation to its delivery

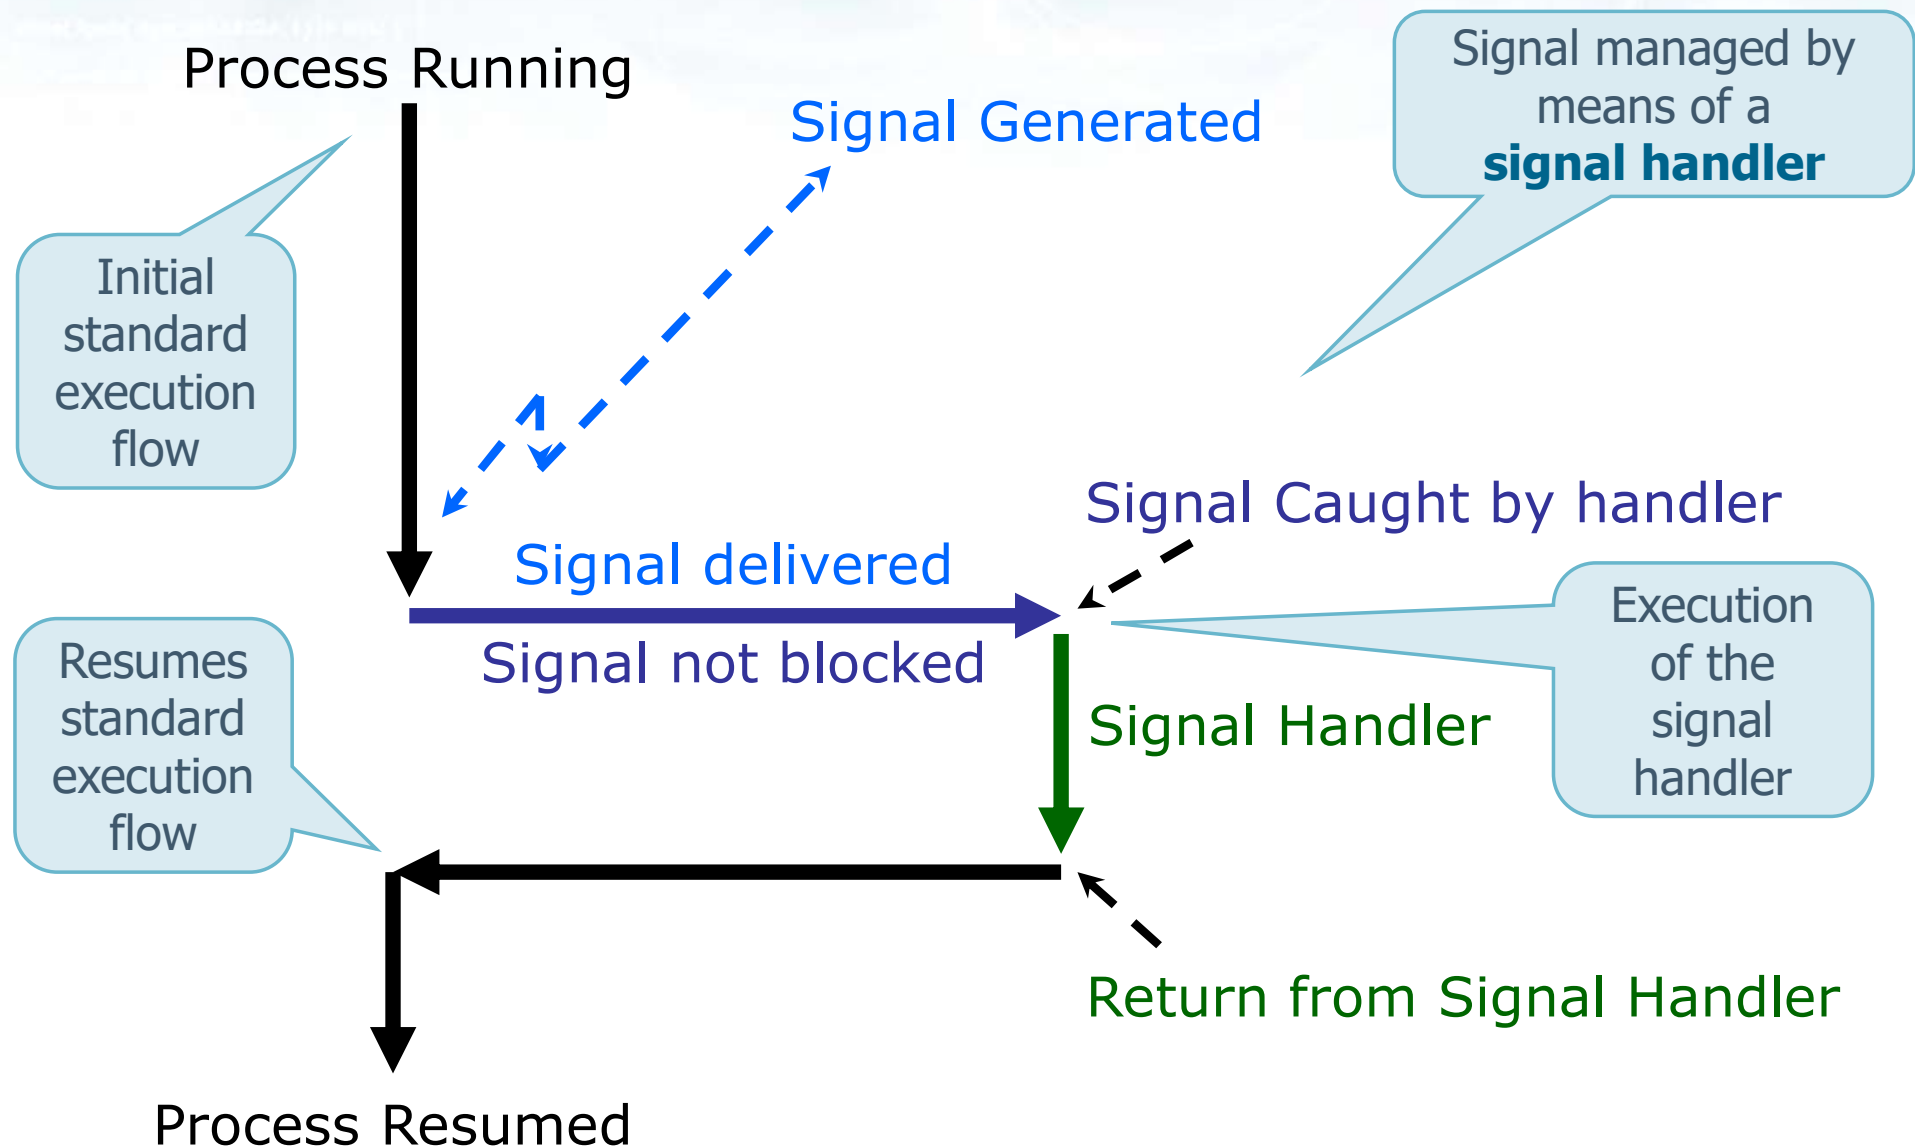
There is no signal queue; the kernel sets a flag in the process table

# Signal management

## ➤ Reaction to a signal

- To properly react to the asynchronous arrival of a given type of signal, a process must inform the kernel about the action that it will perform when it will receive a signal of that type
- A process may
  - **Accept** the **default** behavior (be terminated)
  - Declare to the kernel that it wants to **ignore** the signals of that type
  - Declare to the kernel that it wants to **catch** and manage the signals of that type by means of a signal handler function (similarly to the interrupt management)

# Signal management





# Signal management

❖ Signal management can be carried out with the following system calls

➤ **signal**

- Instantiates a signal handler

➤ **kill (and raise)**

- Sends a signal

The terms **signal** and **kill** are relatively inappropriate. **signal** does not send a signal!!

➤ **pause**

- Suspends a process, waiting the arrive of a signal

➤ **alarm**

- Sends a SIGALARM signal, after a preset time

➤ **sleep**

- Suspends the process for a specified amount of time (waits for signal SIGALRM)



## The signal system call

```
#include <signal.h>
```

```
void (*signal (int sig,  
                void (*func) (int))) (int) ;
```

Received  
parameter  
of the signal  
handler

Returned  
parameter  
of the signal  
handler

- ❖ Allow to instantiate a signal handler
  - Specifies the signal to be managed (**sig**)
  - The function use to manage it (**func**), i.e., the **signal handler**

# The signal system call

```
#include <signal.h>

void (*signal (int sig,
                void (*func) (int))) (int);
```

## ❖ Arguments

- **sig** indicates the type of signal to be caught
  - SIGALRM, SIGUSR1, etc.
- **func** specifies the address (i.e., pointer) to the function that will be executed when a signal of that type is received by the process
  - This function has an argument of `int` type, which indicates the type of signal that will be handled

## The signal system call

```
#include <signal.h>

void (*signal (int sig,
                void (*func) (int))) (int) ;
```

### ❖ Returned values

- on success, the previous value of the signal handler, i.e., the pointer to the previous signal handler function
  - Returns a void \*
- **SIG\_ERR** on error, **errno** is set to indicate the cause
  - #define SIG\_ERR ((void (\*)(int)) -1)

## Reaction to a signal

- ❖ The **signal** system call allows setting three different reactions to the delivery of a signal
  - Accept the default behavior
    - `signal (SIGname, SIG_DFL)`
  - Ignore signal delivery
    - `signal (SIGname, SIG_IGN)`
  - Catch the signal
    - `signal (SIGname, signalHandlerFunction)`

## Example

Setting a program  
to deal with 2  
signals

Same signal handler  
for more than one  
signal type

```
...  
void manager (int sig) {  
    if (sig==SIGUSR1)  
        printf ("Received SIGUSR1\n");  
    else if (sig==SIGUSR2)  
        printf ("Received SIGUSR2\n");  
    else printf ("Received %d\n", sig);  
    return;  
}  
...  
int main () {  
    ...  
    signal (SIGUSR1, manager);  
    signal (SIGUSR2, manager);  
    ...  
}
```

Both signal types  
must be declared

# Example

Asynchronous  
manipulation of  
SIGCHLD (with no wait)

```
static void sigChld (int signo) {  
    if (signo == SIGCHLD)  
        printf ("Received SIGCHLD\n");  
    return;  
}  
...  
signal(SIGCHLD, sigChld);  
if (fork() == 0) {  
    // child  
    ...  
    exit (i);  
} else {  
    // father  
    ...  
}
```

There is no  
pid = wait (&code);

## System call kill

```
#include <signal.h>

int kill (pid_t pid, int sig);
```

- ❖ Send signal (**sig**) to a process or to a group of processes (**pid**)
- ❖ To send a signal to a process, you must have the rights
  - A **user** process can send signals only to processes having the same UID
  - The **superuser** can send signal to any process



# System call kill

## ❖ Arguments

If pid is	Send sig
>0	To process with PID equal to <code>pid</code>
==0	To all processes with GID equal to its GID (if it has the rights)
<0	To all processes with GID equal to the absolute value of <code>pid</code> (if it has the rights)
==−1	To all processes (if it has the rights)

## ❖ Return value

- The value 0, if successful
- The value −1, in case of error

```
int kill (pid_t pid, int sig);
```

## Esercise

- ❖ Write a program using UNIX signals to synchronize two processes  $P_1$  and  $P_2$ 
  - $P_1$  displays a message on standard output and send a message to wake-up  $P_2$
  - $P_2$  wakes-up, displays a message on standard output, and send a message to wake-up  $P_1$
  - This is an example of the expected output
    - Child Woke-up; Father Woke-up; Child Woke-up; Father Woke-up, etc.
- ❖ Use **sleep** instructions, to decrease the probability to have race-conditions

# Solution

```
#include <stdio.h>
#include <unistd.h>
#include <signal.h>
```

```
static void signalHandler ();
static void father (pid_t);
static void child (pid_t);
```

```
static void signalHandler (void) {
    return;
}
```



Signal Handler

# Solution

```
int main (void) {  
    pid_t pid;  
    if (signal(SIGUSR1, signalHandler) == SIG_ERR) {  
        fprintf (stderr, "Signal Handler Error.\n");  
        return (1);  
    }  
    pid = fork ();  
    if (pid==0) {  
        // Child Running  
        child (getppid ());  
    } else {  
        // Father Running  
        father (pid);  
    }  
    return (0);  
}
```

Signal handler  
instantiation

Run process  $P_1$   
and  $P_2$

# Solution

```
static void father (pid_t pid) {  
    fprintf (stdout,  
        "Father Running: pid=%d, childPid=%d\n",  
        getpid (), pid);  
  
    while (1) {  
        sleep (2);  
        kill (pid, SIGUSR1);  
        pause ();  
        fprintf (stdout, "Father Woke-up\n");  
    }  
    return;  
}
```

# Solution

```
static void child (pid_t pid) {  
    fprintf (stdout,  
        "Child Running: pid=%d, fatherPid=%d\n",  
        getpid (), pid);  
    while (1) {  
        pause ();  
        fprintf (stdout, "Child Woke-up\n");  
        sleep (2);  
        kill (pid, SIGUSR1);  
    }  
    return;  
}
```

## System call raise

```
#include <signal.h>

int raise (int sig);
```

- ❖ The **raise** system call allows a process to send a signal to itself
  - The system call
    - `raise (sig)`
  - is equivalent to
    - `Kill (getpid(), sig);`



## System call pause

```
#include <unistd.h>

int pause (void);
```

- ❖ Suspends the calling process until a signal is received
- ❖ Returns after the completion of the signal handler
  - In this case the function returns -1

## System call alarm

```
#include <unistd.h>
```

```
unsigned int alarm (unsigned int seconds);
```

Return value:  
The number of  
seconds remaining or 0


- ❖ Activate a timer (i.e., a count-down)
  - The **seconds** parameter specifies the count-down value (in seconds)
    - At the end of the countdown the signal **SIGALRM** is generated
  - If the system call is executed before the previous call has originated the corresponding signal, the count-down restarts from a new value

## Signal limitations

- ❖ Signals do not convey any information
- ❖ The **memory** of the "pending" signals is **limited**
  - Max one signal pending (sent but not delivered) per type
    - Forthcoming signals of the same type are lost
  - Signals can be ignored
- ❖ Signals require functions that must be **reentrant**
- ❖ Produce **race conditions**
- ❖ Some limitations are avoided in POSIX.4

## Reentrant functions

- ❖ A signal has the following behavior
  - The interruption of the current execution flow
  - The execution of the signal handler
  - The return to the standard execution flow at the end of the signal handler



- ❖ Consequently
  - The kernel **knows** where a signal handler returns, but the signal handler **does not know**
  - The signal handler must operate in a **compatible way** with the original execution flow

## Race conditions

### ❖ Race condition

- The result of more concurrent processes working on common data depends on the execution order of the processes instructions
- ❖ Concurrent programming is subject to race conditions
  - Using signals increases the probability of race conditions
- ❖ Race condition should be avoided at all costs