

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
{
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE *f;
```

```
for(i=0; i<MAXPAROLA; i++)
    freq[i]=0;
```

```
if(argc != 2)
```

```
{
    fprintf(stderr, "ERRORE, serve un parametro con il nome del file\n");
    exit(1);
}
```

```
f = fopen(argv[1], "r");
if(f==NULL)
```

```
{
    fprintf(stderr, "ERRORE, impossibile aprire il file %s\n", argv[1]);
    exit(1);
}
```

```
while( fgets( riga, MAXRIGA, f ) != NULL )
```



# System and Device Programming

## The UNIX I/O

Stefano Quer

Dipartimento di Automatica e Informatica

Politecnico di Torino

## POSIX Standard Library

- ❖ I/O in UNIX can be entirely performed with only **5** functions
  - open, read, write, lseek, close
- ❖ This type of access
  - Is part of POSIX and of the Single UNIX Specification, but not of ISO C
  - It is normally defined with the term "unbuffered I/O", in the sense that each read or write operation corresponds to a system call

## System call `open()`

- ❖ In the UNIX kernel a "file descriptor" is a non-negative integer
- ❖ Conventionally (also for shells)
  - Standard input
    - 0 = `STDIN_FILENO`
  - Standard output
    - 1 = `STDOUT_FILENO`
  - Standard error
    - 2 = `STDERR_FILENO`

These descriptors are defined in the headers file **`unistd.h`**

## System call open()

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int open (const char *path, int flags);

int open (const char *path, int flags,
          mode_t mode);
```

- ❖ It opens a file defining the permissions
- ❖ Return value
  - The descriptor of the file, on success
  - The value -1, on error

## System call `open()`

- ❖ It can have 2 or 3 parameters
  - The **mode** parameter is optional
- ❖ **Path** indicates the file to open
- ❖ **Flags** has multiple options
  - Can be obtained with the OR bit-by-bit of constants defined in the header file **fcntl.h**
  - One of the following three constants is mandatory
    - `O_RDONLY`      open for read-only access
    - `O_WRONLY`      open for write-only access
    - `O_RDWR`        open for read-write access

```
int open (const char *path, int flags, mode_t mode);
```

# System call open

## ➤ Optional constants

- `O_CREAT` creates the files if not exist
- `O_EXCL` error if `O_CREAT` is set and the file exists
- `O_TRUNC` remove the content of the file
- `O_APPEND` append to the file
- `O_SYNC` each write waits that the physical write operation is finished before continuing

...

- `O_NONBLOCK` nonblocking I/O

Analyzed in the next section

```
int open (const char *path, int flags, mode_t mode);
```

# System call open()

❖ **Mode** specifies permission access

- S\_I[RWX]USR      rwx --- ---
- S\_I[RWX]GRP      --- rwx ---
- S\_I[RWX]OTH      --- --- rwx

```
int open (  
    const char *path,  
    int flags,  
    mode_t mode  
);
```

When a file is created, actual permissions are obtained from the **umask** of the user owner of the **process**

## System call read()

```
#include <unistd.h>

int read (int fd, void *buf, size_t nbytes);
```

- ❖ Read from file **fd** a number of bytes equal to **nbytes**, storing them in **buf**
- ❖ Returned values
  - The number of read bytes, on success
  - The value -1, on error
  - The value 0, in the case of EOF



## System call write()

```
#include <unistd.h>
```

```
int write (int fd, void *buf, size_t nbytes);
```

- ❖ Write **nbytes** bytes from **buf** in the file identified by descriptor **fd**
- ❖ Returned values
  - The number of written bytes (i.e., normally **nbytes**), in the case of success
  - The value -1, on error

## System call write()

### ❖ Remarks

- Function **write** writes on the system buffer, not on the disk
  - `fd = open (file, O_WRONLY | O_SYNC);`
- `O_SYNC` forces the sync of the buffers, but only for ext2 file systems

```
int write (int fd, void *buf, size_t nbytes);
```

# Examples

```
float data[10];
```

```
if ( write(fd, data, 10*sizeof(float))==(-1) ) {  
    fprintf (stderr, "Error: Write %d).\n", n);  
}
```

Write a data array (of float values)

```
struct {  
    char name[L];  
    int n;  
    float avg;  
} item;
```

```
if ( write(fd,&item,sizeof(item))==(-1) ) {  
    fprintf (stderr, "Error: Write %d).\n", n);  
}
```

Write a structured item  
(with 3 fields) in binary form

# File Pointers

## ❖ In UNIX

- It is possible to explicitly modify file pointers to perform **random walks** on the file
- Random walks can be implemented using the function `lseek`

## System call lseek()

```
#include <unistd.h>
```

```
off_t lseek (int fd, off_t offset, int whence);
```

- ❖ The current position of the file offset is associated to each file
  - This position indicates the one from which the next read/write operation starts
  - The system call lseek assigns the value **offset** to the file offset

## System call lseek()

- Whence specifies the interpretation of offset
  - If whence==SEEK\_SET
    - The offset is evaluated from the beginning of the file
  - If whence==SEEK\_CUR
    - The offset is evaluated from the current position
  - If whence==SEEK\_END
    - The offset is evaluated from the end of the file

The value of **offset**  
can be positive or  
negative

It is possible to leave  
"holes" in a file  
(filled with zeros)

```
off_t lseek (int fd, off_t offset, int whence);
```

## System call lseek()

- ❖ Return value
  - New offset, on success
  - -1, on error

```
off_t lseek (int fd, off_t offset, int whence);
```

## System call `close()`

```
#include <unistd.h>

int close (int fd);
```

- ❖ It closes the file of descriptor **fd**
  - Notice that, all the open files are closed automatically when the process terminates
- ❖ Return value
  - The value 0, on success
  - The value -1, on error



## Example

```
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>

#define BUFSIZE 4096

int main(void) {
    int nR, nW, fdR, fdW;
    char buf[BUFSIZE];
    fdR = open (argv[1], O_RDONLY);
    fdW = open (argv[2], O_WRONLY | O_CREAT | O_TRUNC,
                S_IRUSR | S_IWUSR);
    if ( fdR==(-1) || fdW==(-1) ) {
        fprintf (stdout, "Error Opening a File.\n");
        exit (1);
    }
}
```

## Example

```
while ( (nR = read (fdR, buf, BUFSIZE)) > 0 ) {  
    nW = write (fdW, buf, nR);  
    if ( nR!=nW )  
        fprintf (stderr,  
            "Error: Read %d, Write %d).\n", nR, nW);  
}  
  
if ( nR < 0 )  
    fprintf (stderr, "Write Error.\n");  
  
close (fdR);  
close (fdW);  
  
exit(0);  
}
```

Error check on the last  
reading operation

This program works indifferently on text and  
binary files