

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
{
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE *f;
```

```
for(i=0; i<MAXPAROLA; i++)
    freq[i]=0;
```

```
if(argc != 2)
```

```
{
    fprintf(stderr, "ERRORE, serve un parametro con il nome del file\n");
    exit(1);
}
```

```
f = fopen(argv[1], "r");
if(f==NULL)
```

```
{
    fprintf(stderr, "ERRORE, impossibile aprire il file %s\n", argv[1]);
    exit(1);
}
```

```
while( fgets( riga, MAXRIGA, f ) != NULL )
```



# System and Device Programming

## File Locking

Stefano Quer

Dipartimento di Automatica e Informatica

Politecnico di Torino

## File locking


- ❖ When two processes edit the same file at the same time the final state of the file corresponds to the last process that wrote the file
- ❖ However, several applications need to be certain that are the only one to write to a file
  - UNIX systems provide file locking
    - It is the term normally used to describe the ability of a process to prevent other processes from modifying a region of a file
  - File locking is a limited form of **synchronization**

## File locking

- ❖ File locking is performed as a byte-range locking
  - A range of a file (possibly, the entire file) is locked
- ❖ For locking we use function `fcntl`
  - The parameter **cmd** must be set to
    - `F_GETLK`, `F_SETLK`, or `F_SETLKW`
  - The third argument must point to an **flock structure**

# The commands

```
int fcntl (int fd, int cmd, struct flock *flockptr);
```



cmd	Purpose
F_GETFL or F_SETFL	Get/set file status flags
F_DUPFD or F_DUPFD_CLOEXEC	Duplicate an existing descriptor
F_GETFD or F_SETFD	Get/set file descriptor flags
F_GETOWN or F_SETOWN	Get/set asynchronous I/O ownership
F_GETLK, F_SETLK, or F_SETLKW	Get/set record locks

Three commands can be issued

# The commands

## ➤ F\_GETLK

- **Check** whether the lock described by **flockptr** is blocked by some other lock
  - If a lock exists that would prevent ours from being created, the information on that existing lock overwrites the information pointed to by flockptr
  - If no lock exists that would prevent ours from being created, the structure pointed to by flockptr is left unchanged except for the l\_type member, which is set to F\_UNLCK

```
int fcntl (int fd, int cmd, struct flock *flockptr);
```

## The commands

### ➤ F\_SETLK

- **Set the lock** described by **flockptr**
- If the compatibility rule prevents the system from giving us the lock `fcntl` returns immediately with `errno` set to either `EACCES` or `EAGAIN`
- This command is also used to clear the lock described by `flockptr` (`l_type` of `F_UNLCK`)

```
int fcntl (int fd, int cmd, struct flock *flockptr);
```

## The commands

### ➤ F\_SETLKW

- **Set the lock** described by **flockptr** with a blocking operation
  - A blocking version of F\_SETLK
  - The W in the command name means wait
- If the requested read lock or write lock cannot be granted because another process currently has some part of the requested region locked, the calling process is put to sleep
- The process wakes up either when the lock becomes available or when interrupted by a signal

```
int fcntl (int fd, int cmd, struct flock *flockptr);
```

# The flock structure

```
int fcntl(int fd, int cmd, struct flock *flockptr);
```

A shared read lock F\_RDLCK  
An exclusive write lock F\_WRLCK  
Unlocking a region F\_UNLCK

Offset in bytes  
Relative to l\_whence

Length, in bytes  
0 means lock to EOF

```
struct flock {  
    short l_type;  
    short l_whence;  
    off_t l_start;  
    off_t l_len;  
    pid_t l_pid;  
};
```

SEEK\_SET,  
SEEK\_CUR, or  
SEEK\_END

Returned with F\_GETLK



# The flock structure

```
int fcntl(int fd, int cmd, struct flock *flockptr);
```

The starting byte offset  
of the region being  
locked or unlocked  
(l\_start and l\_whence)

SEEK\_SET,  
SEEK\_CUR, or  
SEEK\_END

```
struct flock {  
    short l_type;  
    short l_whence;  
    off_t l_start;  
    off_t l_len;  
    pid_t l_pid;  
};
```

The size of  
the region in  
bytes (l\_len)

hFile

Locked portion

Returned with command  
F\_GETLK

# The flock structure

```
int fcntl(int fd, int cmd, struct flock *flockptr);
```

The ID (l\_pid) of the process holding the lock that can block the current process (returned by F\_GETLK only)

```
struct flock {  
    short l_type;  
    short l_whence;  
    off_t l_start;  
    off_t l_len;  
    pid_t l_pid;  
};
```

SEEK\_SET,  
SEEK\_CUR, or  
SEEK\_END

Returned with command  
F\_GETLK

## Guidelines

### ❖ Several rules apply

- The two parameters specifying the starting offset of the region are similar to the last two arguments of the **lseek** function

```
off_t lseek (int fd, off_t offset, int whence);
```

- Locks can start and extend beyond the current end of file, but cannot start or extend before the beginning of the file

```
struct flock {  
    short l_type;  
    short l_whence;  
    off_t l_start;  
    off_t l_len;  
    pid_t l_pid;  
};
```

## Guidelines

- If **l\_len** is 0, the lock extends to the largest possible offset of the file
  - This allows us to lock a region starting anywhere in the file, up through and including any data that is appended to the file
- To lock the entire file, we set **l\_start** and **l\_whence** to point to the beginning of the file and specify a length (**l\_len**) of 0

```
struct flock {  
    short l_type;  
    short l_whence;  
    off_t l_start;  
    off_t l_len;  
    pid_t l_pid;  
};
```

# Guidelines

## ➤ Repeated Lock Request

- If a lock is present
- When a new lock request is granted or refused ?


Exiting Lock	Requested Lock Type	
	Read Lock	Write Lock
None	Granted	Granted
Read lock	Granted	Refused
Write lock	Refused	Refused

## Guidelines

- Lock **belongs** to a **process**, and it is possible to
  - The compatibility rule applies to lock requests made from different processes, not to multiple lock requests made by a single process
  - If a process has an existing lock on a range of a file, a subsequent attempt to place a lock on the same range by the same process will replace the existing lock with the new one
    - If a process has a write lock on bytes 16–32 of a file and then tries to place a read lock on bytes 16–32, the request will succeed, and the write lock will be replaced by a read lock

# Guidelines

- When setting or releasing a lock on a file, the system combines or splits adjacent areas as required



## Guidelines

### ➤ File locking can produce

- Starvation
  - Processes A and B periodically obtain a shared lock whereas C is waiting forever for an exclusive lock
- Deadlock
  - Process A is waiting for B to unlock and vice-versa (even on slightly a different file region)



# Example

## Lock a file region

```
#include <fcntl.h>
```

```
...
```

```
int lock_region (int fd, int cmd, int type, off_t offset,  
                int whence, off_t len) {
```

```
    struct flock lock;
```

```
    lock.l_type = type;  
    lock.l_start = offset;  
    lock.l_whence = whence;  
    lock.l_len = len;
```

```
    return (fcntl(fd, cmd, &lock));
```

```
}
```

Offset in bytes, relative to l\_whence

A shared read lock F\_RDLCK  
An exclusive write lock F\_WRLCK  
Unlocking a region F\_UNLCK

SEEK\_SET, SEEK\_CUR, or SEEK\_END

Length, in bytes  
0 means lock to EOF

-1 on error