

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
{
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE *f;
```

```
for(i=0; i<MAXPAROLA; i++)
    freq[i]=0;
```

```
if(argc != 2)
```

```
{
    printf(stderr, "ERRORE, serve un parametro con il nome del file\n");
    exit(1);
}
```

```
f = fopen(argv[1], "r");
if(f==NULL)
```

```
{
    printf(stderr, "ERRORE, impossibile aprire il file %s\n", argv[1]);
    exit(1);
}
```

```
while( fgets( riga, MAXRIGA, f ) != NULL )
```



System and Device Programming

I/O Multiplexing


Stefano Quer

Dipartimento di Automatica e Informatica

Politecnico di Torino

I/O Multiplexing

- ❖ Let us suppose we must read from one descriptor and write to another




- We can use blocking I/O in a loop

```
while ((n = read(STDIN_FILENO, buf, BUFSIZ)) > 0)
    if (write(STDOUT_FILENO, buf, n) != n)
        err_sys("write error");
```

- This form of blocking I/O is very common

I/O Multiplexing

❖ How can we read from **two** descriptors?



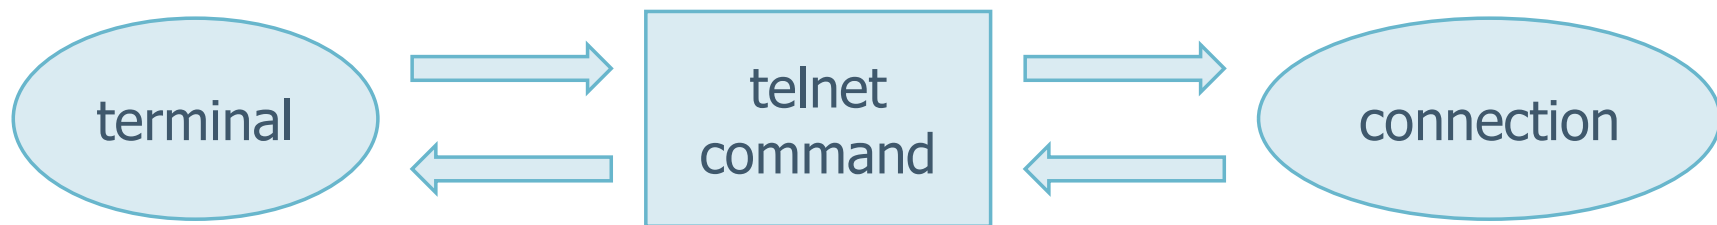
➤ We cannot do a blocking read on either descriptor

- Data may appear on **one** descriptor while we are blocked in a read on the **other**
- We need a different technique to handle this case and we have several possibilities
 - Use different processes or different threads
 - Use non-blocking I/O
 - Use asynchronous I/O
 - Use I/O multiplexing

Example

❖ A telnet command


- Reads from the terminal (standard input) and write to a network connection
- Reads from the network connection and write to the terminal (standard output)



- The telnet process has two inputs and two outputs
 - We can't do a blocking read on either of the inputs, as we never know which input will have data for us

Solution 1

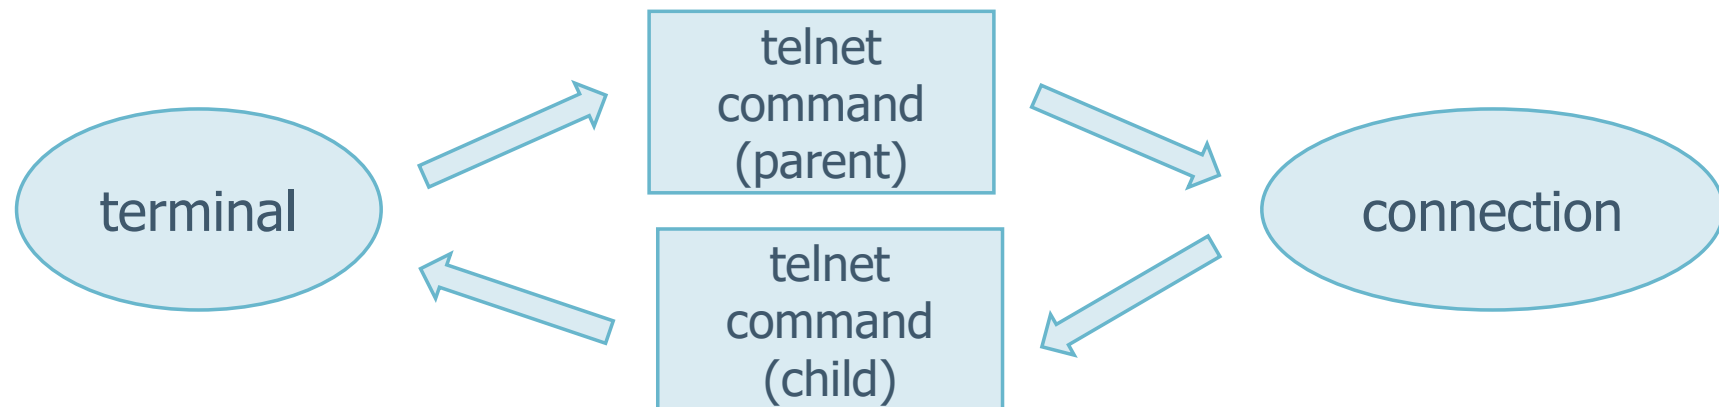
- ❖ The multi-process approach to multiplexing
 - Fork the process, with the parent handling one direction of data and the child the other one



- Each process can perform blocking read
- When one process receives and EOF must stop the other one (e.g., sending a signal)

Solution 2

- ❖ The multi-threads approach to multiplexing
 - Same as before but two threads instead of two processes



- This avoids the termination complexity
- It requires that we deal with synchronization between the threads, which could add more complexity than it saves

Solution 3

- ❖ The non-blocking I/O approach to multiplexing
 - We use nonblocking I/O in a single process
 - We set both descriptors to be nonblocking
 - We issue a read on the first descriptor
 - If data is present, we read it and process it
 - If there is no data to read, the call returns and we do the same on the second descriptor
 - We wait for some amount of time and then we try to read from the first descriptor again
 - This type of loop is called **polling**
 - Most of the time, there won't be data to read, so we waste time performing the read system calls
 - We have to guess how long to wait before looping

Solution 4

Please, see Section u04s08

- ❖ The asynchronous I/O approach to multiplexing
 - We tell the kernel to notify us with a signal when a descriptor is ready for I/O
 - This approach has two problems
 - Portability can be an issue
 - Different systems provide their own limited forms of asynchronous
 - POSIX proposes one strategy
 - System V provides the SIGPOLL
 - BSD has SIGIO
 - etc.

Solution 4

- The number of signals allowed for each process can be a problem
 - A few implementation of asynchronous I/O allow only one signal
 - POSIX allows us to select which signal to use for the notification
 - Anyway, the number of signals we can use is still far less than the number of possible open file descriptors
 - To determine which descriptor is ready, we would need to set each file descriptor to nonblocking mode and try the descriptors in sequence
 - This is again, a form of **polling**

Solution 5

- ❖ The “I/O multiplexing” approach to multiplexing
 - The idea is to use a function that
 - Accepts more than one I/O descriptor
 - Doesn't return until one of the descriptors is ready for I/O
 - On return, it tells us which descriptors are ready for I/O

Solution 5

❖ The “I/O multiplexing” logic flow

➤ We build a list of the descriptors that we are interested in

- Usually more than one descriptor
- We need to perform I/O on each one of them

➤ We use a multiplexing I/O function

- Three functions allow us to perform I/O multiplexing
 - select, pselect, poll

Functions pselect and poll are variation of select

- They differ from the user interface point of view
- Include POSIX `<sys/select.h>`
 - Or `<sys/types.h>`, `<sys/time.h>`, and `<unistd.h>`

The select function

```
#include <sys/select.h>
int select (
    int maxfdp1,
    fd_set *restrict readfds,
    fd_set *restrict writefds,
    fd_set *restrict exceptfds,
    struct timeval *restrict tvptr
);
```

- ❖ Function **select** allows I/O multiplexing under all POSIX-compatible platforms
- ❖ The arguments tell the kernel
 - Which descriptors we are interested in
 - Which conditions we focus on (for each descriptor)
 - How long we want to wait

The select function

❖ The last argument specifies how long we want to wait in terms of seconds and microseconds

➤ If `tvptr == NULL`

- We wait forever
- If a signal interrupts the function, select returns -1 with `errno` set to `EINTR`

➤ If `tvptr->tv_sec == 0 && tvptr->tv_usec == 0`

- We do not wait at all
- All the specified descriptors are tested, and return is made immediately
- This is a way to poll

```
int select (  
    int maxfdp1,  
    fd_set *restrict readfds,  
    fd_set *restrict writefds,  
    fd_set *restrict exceptfds,  
    struct timeval *restrict tvptr  
);
```

The select function

- If `tvptr->tv_sec != 0 || tvptr->tv_usec != 0`
 - We wait the specified number of seconds and microseconds
 - Return is made when one of the specified descriptors is ready or when the timeout value expires
 - If the timeout expires before any of the descriptors is ready, the return value is 0

```
int select (  
    int maxfdp1,  
    fd_set *restrict readfds,  
    fd_set *restrict writefds,  
    fd_set *restrict exceptfds,  
    struct timeval *restrict tvptr  
);
```

The select function

- ❖ The middle three arguments are pointers to descriptor sets
 - They specify the descriptors we check and for which conditions
 - Readable, writable, or exceptions
 - Any of them can be null pointers if we are not interested
 - A descriptor set is stored in an **fd_set** data type
 - We can consider it to be just an array of bits
 - If all three pointers are NULL, then we have a higher-precision timer (provided by sleep)

```
int select (  
    int maxfdp1,  
    fd_set *restrict readfds,  
    fd_set *restrict writefds,  
    fd_set *restrict exceptfds,  
    struct timeval *restrict tvptr  
);
```


The select function

- Objects of type `fd_set` data type can be manipulated by a function (or macro) to
 - Set to all zero bits, by `FD_ZERO`
 - Turn on a single bit in a set, by `FD_SET`
 - Clear a single bit, by `FD_CLR`
 - Test whether a given bit is turned on, by `FD_ISSET`

```
#include <sys/select.h>
```

```
void FD_SET (int fd, fd_set *fdset);  
void FD_CLR (int fd, fd_set *fdset);  
void FD_ZERO (fd_set *fdset);  
int FD_ISSET (int fd, fd_set *fdset);
```

Example

- ❖ The following code show how to prepare a read set

```
fd_set rset;  
int fd;
```

We define the set

```
FD_ZERO (&rset);  
FD_SET (fd, &rset);  
FD_SET (STDIN_FILENO, &rset);
```

We zero the set

We set bits for the descriptor we are interested in

- ❖ Once, select returns, we tests whether a given bit is still set

```
if (FD_ISSET(fd, &rset)) {  
    ...  
}
```

The select function

- ❖ The first argument to select stands for
 - The “maximum file descriptor plus 1.”
 - We compute the highest descriptor that we are interested in, considering all three of the descriptor sets, add 1
 - We could just set the first argument to FD_SETSIZE, a constant in <sys/select.h> that specifies the maximum number of descriptors (often 1,024), but this value is too large for most applications

```
int select (  
    int maxfdp1,  
    fd_set *restrict readfds,  
    fd_set *restrict writefds,  
    fd_set *restrict exceptfds,  
    struct timeval *restrict tvptr  
);
```

The select function

- ❖ There are three possible return values
 - A value of -1 means that an error occurred
 - A return value of 0 means that no descriptors are ready
 - This happens if the time limit expires before any of the descriptors are ready
 - When this happens, all the descriptor sets will be zeroed out

```
int select (  
    int maxfdp1,  
    fd_set *restrict readfds,  
    fd_set *restrict writefds,  
    fd_set *restrict exceptfds,  
    struct timeval *restrict tvptr  
);
```

The select function

- A positive return value specifies the number of descriptors that are ready
 - This value is the sum of the number of descriptors ready in all three sets
 - If the same descriptor is ready to be read and written, it will be counted twice in the return value
 - The only bits left on in the three descriptor sets are the bits corresponding to the descriptors that are ready

```
int select (  
    int maxfdp1,  
    fd_set *restrict readfds,  
    fd_set *restrict writefds,  
    fd_set *restrict exceptfds,  
    struct timeval *restrict tvp  
) ;
```

The select function

- ❖ A descriptor is consider as “ready”
 - If a read from that descriptor won't block (if it is in the read set, readfds)
 - If a write from that descriptor won't block (if it is in the write set, writefds)
 - If an exception condition is pending on that descriptor (if it is in the exception set, exceptfds)
 - File descriptors for regular files always return ready for reading, writing, and exception conditions

```
int select (  
    int maxfdp1,  
    fd_set *restrict readfds,  
    fd_set *restrict writefds,  
    fd_set *restrict exceptfds,  
    struct timeval *restrict tvptr  
);
```