

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
{
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE *f;
```

```
for(i=0; i<MAXPAROLA; i++)
    freq[i]=0;
```

```
if(argc != 2)
```

```
{
    fprintf(stderr, "ERRORE, serve un parametro con il nome del file\n");
    exit(1);
}
```

```
f = fopen(argv[1], "r");
if(f==NULL)
```

```
{
    fprintf(stderr, "ERRORE, impossibile aprire il file %s\n", argv[1]);
    exit(1);
}
```

```
while( fgets( riga, MAXRIGA, f ) != NULL )
```



# System and Device Programming

## Asynchronous I/O

Stefano Quer

Dipartimento di Automatica e Informatica

Politecnico di Torino

## Synchronous I/O

- ❖ All standard I/O operations are synchronous
  - I/O is **blocking** and the task waits until the I/O operation completes
  - Select delivers a synchronous form of notification
- ❖ Unfortunately, synchronous operations are inherently **slow** compared to other processing
  - Delays may be caused by
    - Hardware device, e.g., track and sector seek time on random access, etc.
    - Relatively slow data transfer rate between a physical device and the system memory
    - Network transfer using file servers, storage area networks, etc.

## Asynchronous I/O

- ❖ Threads can perform asynchronous I/O
  - A task can continue **without** waiting for an I/O operation to complete
- ❖ There are two techniques
  - Multithread I/O
  - POSIX asynchronous I/O

# Asynchronous I/O

## ❖ Multithread I/O

- We use multiple threads
- Each thread within a process (or in different processes) may use a synchronous model
  - Each thread is responsible for a sequence of one or more synchronous, **blocking** I/O operations
  - Each thread should have its own file or pipe handle
- **Other** threads can **continue** execution
  - The threads run asynchronous to each other
- This is the most general technique

# Asynchronous I/O

## ❖ Asynchronous I/O

### ➤ We incur additional complexity when we use the POSIX asynchronous I/O interfaces

- We have to worry about sources of errors for every asynchronous operation
- The interfaces involve a lot of extra setup and processing rules compared to their conventional counterparts
- Recovering from errors can be difficult
  - For example, if we submit multiple asynchronous writes and one fails, how should we proceed? If the writes are related, we might have to undo the ones that succeeded

## The aiocb data structure

- ❖ POSIX gives us a consistent way to perform asynchronous I/O, regardless of the type of file
- ❖ The interfaces use AIO control blocks to describe I/O operations
  - The aiocb structure defines an AIO control block
  - It contains at least the fields shown in the following structure (implementations might include additional fields)
  - Note
    - Insert the `<aio.h>` library in the C file
    - Compile the C file with the realtime library (librt.a), i.e., `-lrt`

# The aiocb data structure

```
#include <aio.h>

struct aiocb {
    int aio_fildes;
    off_t aio_offset;
    volatile void *aio_buf;
    size_t aio_nbytes;
    int aio_reqprio;
    struct sigevent aio_sigevent;
    int aio_lio_opcode;
};
```

## ❖ Parameters

- The **aio\_fildes** field is the file descriptor open for the file to be read or written



## The aiocb data structure

- Read or writes start at the offset specified by **aio\_offset**
  - For a read Data is copied to the buffer that begins at the address specified by `aio_buf`
  - For a write, data is copied from this buffer
- The **aio\_nbytes** field contains the number of bytes to read or write

```
struct aiocb {  
    int aio_fildes;  
    off_t aio_offset;  
    volatile void *aio_buf;  
    size_t aio_nbytes;  
    int aio_reqprio;  
    struct sigevent aio_sigevent;  
    int aio_lio_opcode;  
};
```



## The aiocb data structure

- The **aio\_reqprio** field is a hint that gives applications a way to suggest an ordering for the asynchronous I/O requests
  - The system has only limited control over the exact ordering, however, so there is no guarantee that the hint will be honored

```
struct aiocb {
    int aio_fildes;
    off_t aio_offset;
    volatile void *aio_buf;
    size_t aio_nbytes;
    int aio_reqprio;
    struct sigevent aio_sigevent;
    int aio_lio_opcode;
};
```

## The aiocb data structure

- The **aio\_sigevent** field controls how the application is notified about the completion of the I/O event
  - It is described by a sigevent structure

```
struct sigevent {  
    int sigev_notify;  
    int sigev_signo;  
    union sigval sigev_value;  
    void (*sigev_notify_function)(union sigval);  
    pthread_attr_t *sigev_notify_attributes;  
};
```

## The aiocb data structure

- The `sigev_notify` field controls the type of notification
  - `SIGEV_NONE` the process is not notified when the asynchronous I/O request completes
  - `SIGEV_SIGNAL` the signal specified by the `sigev_signo` field is generated when the asynchronous I/O request completes
  - `SIGEV_THREAD` the function specified by the `sigev_notify_function` field is called when the asynchronous I/O request completes  
address of a pthread attribute

```
struct sigevent {  
    int sigev_notify;  
    int sigev_signo;  
    union sigval sigev_value;  
    void (*sigev_notify_function)(union sigval);  
    pthread_attr_t *sigev_notify_attributes;  
};
```

## The aiocb data structure

- The **aio\_reqprio** field is a hint that gives applications a way to suggest an ordering for the asynchronous I/O requests
  - The system has only limited control over the exact ordering, however, so there is no guarantee that the hint will be honored

```
struct aiocb {
    int aio_fildes;
    off_t aio_offset;
    volatile void *aio_buf;
    size_t aio_nbytes;
    int aio_reqprio;
    struct sigevent aio_sigevent;
    int aio_lio_opcode;
};
```

## Read and Write

```
#include <aio.h>
```

```
int aio_read(struct aiocb *aiocb);
```

```
int aio_write(struct aiocb *aiocb);
```

- ❖ To perform asynchronous I/O, we need to
  - Initialize an AIO control block
  - Call either **aio\_read** or the **aio\_write**
- ❖ When these functions return success, the asynchronous I/O request has been queued for processing by the operating system
  - The return value has no relation to the result of the actual I/O operation

## Guidelines

- ❖ While the I/O operation is pending, we have to be careful to ensure that the AIO control block and data buffer remain stable
  - Their underlying memory must remain valid and we cannot reuse them until the I/O operation completes

```
int aio_read(struct aiocb *aiocb);  
int aio_write(struct aiocb *aiocb);
```

## Synchronization

```
#include <aio.h>
```

```
int aio_fsync (int op, struct aiocb *aiocb);
```

- ❖ We can use this function to force all pending asynchronous writes to persistent storage without waiting
- ❖ The `aio_fsync` operation returns when the synch is scheduled
  - The data will not be persistent until the asynchronous synch completes
  - The AIO control block controls how we are notified



## Synchronization

- ❖ The `aio_fildes` field in the AIO control block indicates the file whose asynchronous writes are synched
- ❖ If the `op` argument is set to
  - `O_DSYNC`, then the operation behaves like a call to `fdatasync`
  - `O_SYNC`, the operation behaves like a call to `fsync`

```
int aio_fsync (int op, struct aiocb *aiocb);
```

## Suspension

```
#include <aio.h>
```

```
int aio_suspend(const struct aiocb *const list[],  
               int nent, const struct timespec *timeout);
```

- ❖ We use asynchronous I/O when we have other processing to do and we don't want to block while performing the I/O operation
- ❖ However, when we have completed the processing and find that we still have asynchronous operations outstanding, we can call the `aio_suspend` function to block until an operation completes

**Cancel**

```
#include <aio.h>
```

```
int aio_cancel (int fd, struct aiocb *aiocb);
```

- ❖ When we have pending asynchronous I/O operations that we no longer want to complete, we can attempt to cancel them with the `aio_cancel` function