

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
{
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE *f;
```

```
for(i=0; i<MAXPAROLA; i++)
    freq[i]=0;
```

```
if(argc != 2)
```

```
{
    fprintf(stderr, "ERRORE, serve un parametro con il nome del file\n");
    exit(1);
}
```

```
f = fopen(argv[1], "r");
if(f==NULL)
```

```
{
    fprintf(stderr, "ERRORE, impossibile aprire il file %s\n", argv[1]);
    exit(1);
}
```

```
while( fgets( riga, MAXRIGA, f ) != NULL )
```



System and Device Programming

Advanced UNIX IPC

Stefano Quer

Dipartimento di Automatica e Informatica

Politecnico di Torino

Pipes

- ❖ Original pipes (or **unnamed** pipes) are a common form of UNIX System IPC but they
 - Can be used only between two processes that have a **common ancestor**
 - Last only as long as the process last
 - Are **half duplex**
 - Data flows only in one direction
 - Can be seen as an **unstructured** sequential files
 - The communication channel is a FIFO queue
 - Structured data transfers require a communication protocol
 - Are limited to transfer only **limited quantity** of memory

Communication channels

❖ This section extends pipes in the following directions

➤ Pipes extensions

- FIFOs
- Message queues

➤ Shared memory

To avoid a common ancestor

To allow structured data

To transfer a lot on information

Memory mapped files have been introduced in the filesystem unit

FIFOs

- ❖ FIFOs are an extension of traditional pipes and are sometimes called **named pipes**
 - They allow a communication among **unrelated** processes
 - They can last as long as the system does
 - They can be deleted if no longer used
- ❖ A FIFO is a type of file
 - Creating a FIFO is similar to creating a file
 - Indeed, a FIFO corresponds to a file in the local storage
 - They have a **pathname** in the filesystem
 - Once a FIFO has been created, processes can open it and perform R/W operations on it

FIFOs

❖ Process logic

- A FIFO special file is entered into the filesystem by calling **mkfifo**
 - **Subsequent** calls to (the “same”) **mkfifo** **have no effect**
- Once we have created a FIFO special file, any process can open it, using the **open** system call
- Once a FIFO has been opened, it can be used for reading or writing
 - Use **read** and **write**, as for ordinary files
- Notice that a FIFO has to be open at **both** ends before you can proceed to do any input or output operations on it

FIFOs

```
#include <sys/stat.h>
```

```
int mkfifo (const char *path, mode_t mode);
```

❖ Function **mkfifo** creates a FIFO

➤ The parameters **path** and **mode** are similar to the corresponding ones specified for function **open**

- Please refer to **open** for any further explanation on the **mode** parameter
 - Use constant `S_I[RWX]USR` or an octal representation to specify user and group ownership

FIFOs

- ❖ Return value
 - The value 0, on success
 - The value -1, on error
- ❖ Once the FIFO is in the system, we can use normal file I/O functions to operate on it
 - Please refer to system calls `open`, `read`, `write`, and `close` for further details

It creates a FIFO with pathname, e.g.,
prw-rw-r-- 1 quer quer 0 apr 5 16:14 **path**

```
int mkfifo (const char *path, mode_t mode);
```


FIFOs

```
#include <sys/stat.h>
```

```
int mkfifoat (int fd, const char *path, mode_t mode);
```

- ❖ Function **mkfifoat** is similar to **mkfifo** but
 - Parameter **fd** indicates a directory path to use to open the FIFO file
 - If **path** specifies an absolute pathname, **fd** is ignored (and **mkfifoat** is equivalent to **mkfifo**)
 - If **path** specifies a relative pathname and **fd** is a descriptor for an open directory, the pathname is evaluated starting from this directory
 - If **path** specifies a relative pathname and **fd** is **AT_FDCWD** the pathname is evaluated starting from the current working directory

FIFOs

❖ Caveats

- As with a pipe, if we write to a FIFO that no process has opened for reading, the signal SIGPIPE is generated
- When the last writer for a FIFO closes the FIFO, an end of file is generated for the reader of the FIFO
- It is common to have multiple writers for a given FIFO
 - We have to worry about atomic writes if we don't want the outputs from multiple processes to be interleaved

Example

Run this process on
a shell windows

1 Reader + 1 Writer
(client server communication)

```
int main() {
    int fd; char str[80];
    char *myfifo = "/tmp/myfifo";
    mkfifo (myfifo, 0666);
    fd = open (myfifo, O_WRONLY);
    while (1) {
        printf ("Send to reader: ");
        fgets (str, 80, stdin);
        write (fd, str, strlen (str)+1);
        if (strncmp (str, "end", 3)==0) {
            break;
        }
    }
    close (fd);
    return 0;
}
```

(W) P₁

(R) P₂

The Writer

Read from stdin

Write to FIFO

Stop the process
when "end" is
introduced

Example

Run this process on
another shell windows

```
int main() {  
    int fd;  
    char str[80];  
    char *myfifo = "/tmp/myfifo";  
    mkfifo (myfifo, 0666);  
    fd1 = open (myfifo, O_RDONLY);  
    while (1) {  
        read (fd, str, 80);  
        printf ("Received from writer: %s", str);  
        if (strncmp (str, "end", 3)==0) {  
            break;  
        }  
    }  
    close(fd);  
    return 0;  
}
```

(W) P₁

(R) P₂

The Reader

Read from the FIFO

Stop the process
when "end" is
received

Blocking versus Non-blocking

- ❖ The **open** operation on a FIFO can be blocking or non-blocking

```
fd = open (myfifo, ... | O_NONBLOCK);
```

- Without the **O_NONBLOCK** flag
 - On **open** in read-only (write-only) mode is **blocking** until some other process open the FIFO in write-only (read-only)
- With the **O_NONBLOCK** flag
 - An open in read-only mode return immediately
 - An open in write-only mode returns -1 (and errno set to ENXIO)

Error Checking

- ❖ Many UNIX system functions (such as **mkfifo** and **mkfifoat**) returns -1 on error
- ❖ Once this happens, the strategy to check the origin of the error is the following one
 - Include header **<errno.h>**
 - This header defines error codes and error manipulation functions
 - Define the global integer variable **errno**
 - This variable is automatically set to the proper error
 - Use functions **perror** and **strerror** to display an error message

Error Checking

When **mkfifo** generates an error, the following error codes are set in **errno**

```
int mkfifo (path, S_IRWXU | S_IRWXG | S_IRWXO);
```

Error Code	Error Meaning
EACCES	One of the directories in path did not allow search/execute permission
EDQUOT	The user's quota on the filesystem has been exhausted
EEXIST	path already exists This includes the case where path is a symbolic link, dangling or not
ENAMETOOLONG	Either the total length of path is greater than PATH_MAX , or an individual filename component has a length greater than NAME_MAX
ENOENT	A directory component in path does not exist or is a dangling symbolic link
ENOSPC	The directory or filesystem has no room for the new file
ENOTDIR	A component used as a directory in path is not, in fact, a directory
EROFS	path refers to a read-only filesystem

Example

To the Reader and the Writer
analyzed before **add**

Define the
errno header

```
#include <errno.h>
```

```
extern int errno;
```

```
int ret;
```

```
ret = mkfifo (myfifo, 0666);
```

```
if (errno == EEXIST)  
    fprintf (stderr, "FIFO exists.\n");
```

```
sprintf (str, "Reader (return value=%d)", ret);  
perror (str);
```

Define the **errno** variable
(automatically set by the system in case
of an error, for many system calls)

Grab error code (-1)
from mkfifo

Manually/Explicitly
check for error

Use **perror** to display
error condition

Function **perror** displays the string str
followed by ":", an error message, and "\n"

Example

❖ Possible error messages

```
sprintf (str, "... (return value=%d)", ret);  
error (str);
```

Writer (return value=0): **Success**

Reader (return value=-1): **File exists**


❖ FIFO may be eventually removed

```
sprintf (str, "rm -rf %s", myfifo);  
system (str);
```


To remove "/tmp/myfifo", i.e.,
prw-rw-r-- 1 quer quer 0 apr 5 16:14 myfifo

Example: Extensions

- ❖ Alternate Write and Read operations
 - Use same FIFO in both directions



- ❖ Coordinate more Writers with one Reader
 - Use more FIFOs



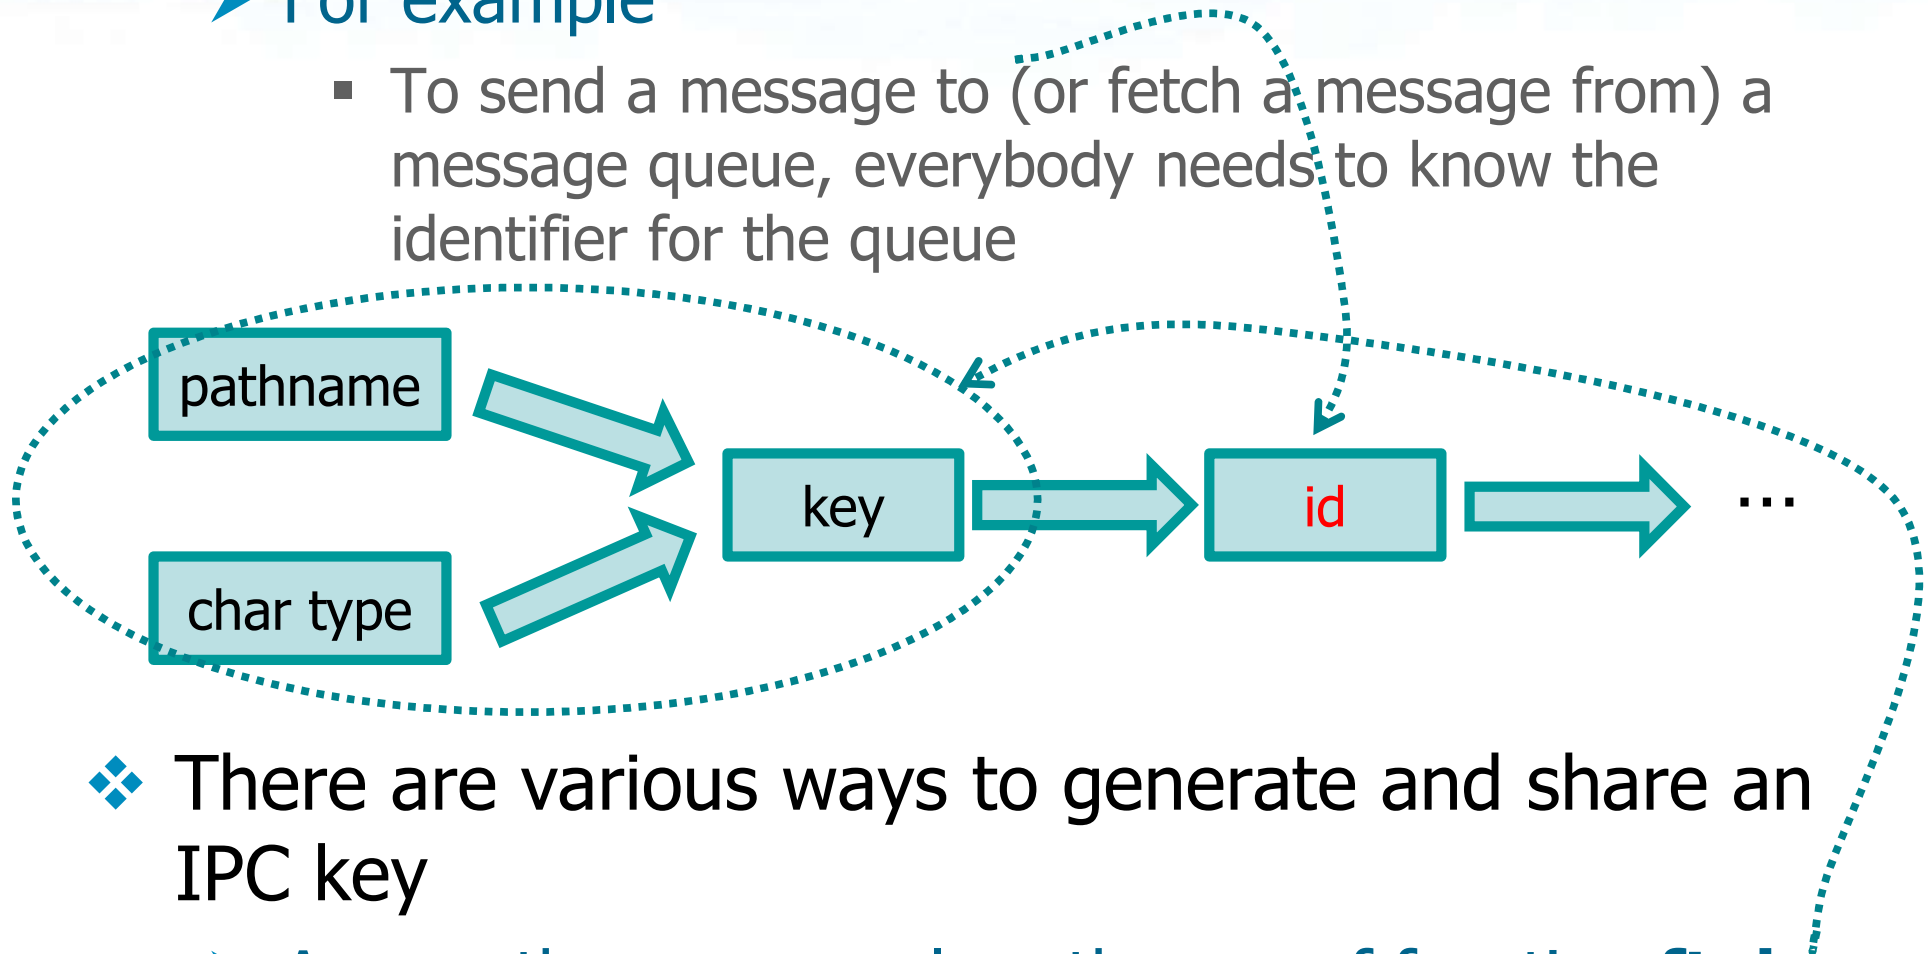
Identifiers and keys

- ❖ Message queues and shared memories are related to IPC structures
 - Whenever an IPC structure is being created a key must be specified
 - The data type of this key is the primitive system datatype **key_t**
 - **key_t** is often defined as a long integer in the header `<sys/types.h>`
 - The kernel converts this key into an identifier
 - Each IPC structure is referred in the kernel by a non-negative integer identifier
 - This identifier is like an internal reference to the object

Identifiers and keys

➤ For example

- To send a message to (or fetch a message from) a message queue, everybody needs to know the identifier for the queue



- ❖ There are various ways to generate and share an IPC key

➤ Among them, we analyze the use of function **ftok**

Identifiers and keys

```
#include <sys/msg.h>

key_t ftok (const char *path, int id);
```

- ❖ To share a key, the different processes (e.g., the clients and the server) must agree on a
 - Standard (file) pathname
 - The file **must** exist
 - Project ID
 - A character value between 0 and 255
- ❖ Function **ftok** converts these two parameters into a key (of type **key_t**)

Identifiers and keys

- ❖ This key is then used during the communication phase
 - Notice that the only service provided by **ftok** is a way of generating a key from a pathname and a project ID
 - Function **ftok** does not perform any sort of communication
- ❖ Return value
 - Message key, if success
 - The value -1 , on error

```
key_t ftok(const char *path, int id);
```

Message queues

- ❖ FIFOs are used to pass streams of anonymous bytes
 - Applications using FIFOs have to manage their own data chunking
 - They have to agree on data delimiters, such as end-of-field, end-of-record, etc.
- ❖ To pass structured data chunks it is necessary to use message queues
 - A message queue is a linked list of messages stored within the kernel and identified by a message queue identifier

Message queues

What is a queue?

❖ A message queue

- Is created or an existing queue opened by **msgget**
- The queue may be controlled using **msgctl**
- New messages are added to the end of a queue by **msgsnd**
- Messages are fetched from a queue by **msgrcv**
 - Messages do not have to be fetched in a first-in, first-out order
 - Messages can be fetched based on their type field

All right ...
A message queue manipulates messages.
But what is a message?

Message queues

What is a message?

❖ Each message manipulated by `msgsnd`

➤ Is composed of

- A positive long integer type field
- A non-negative length (`N_BYTES`)
- The actual data bytes (`mtext`) of size `nbytes`

➤ Has to be defined by the user as a C data structure including

- The message type `mtype`
- The data field `mtext`
of size `N_BYTES`

```
struct mymesg {  
    long int mtype;  
    char mtext[N_BYTES];  
};
```

➤ Messages are always placed at the end of the queue

Message get

```
#include <sys/msg.h>

int msgget (key_t key, int flag);
```

- ❖ Function **msgget** either open an existing queue or create a new queue
 - Key is the values generated with **ftok**
 - Flag is used to define the mode permission field to a data structure associated to the message queue
- ❖ Return value
 - Message queue identifier (**msqid**), if success
 - The value **-1**, on error

Message control

```
#include <sys/msg.h>

int msgctl (
    int msqid, int cmd, struct msqid_ds *buf
);
```

- ❖ Function **msgctl** performs various operations on a queue
 - The queue is specified by its identifier (msqid)
 - The parameter msqid is the value returned by msgget

Message control

- The `cmd` argument specifies the command to be performed on the queue
 - `IPC_STAT`
 - Fetch the `msqid_ds` structure for this queue, storing it in the structure pointed to by `buf`
 - `IPC_SET`
 - Copy the following fields from the structure pointed to by `buf` to the `msqid_ds` structure associated with this queue: `msg_perm.uid`, `msg_perm.gid`, `msg_perm.mode`, and `msg_qbytes`

```
int msgctl (int msqid, int cmd, struct msqid_ds *buf);
```

Message control

- **IPC_RMID**
 - Remove the message queue from the system and any data still on the queue
 - The removal is immediate
 - Any other process still using the message queue will get an error of EIDRM on its next attempted operation on the queue
- These three constants are also provided for semaphores and shared memory

```
int msgctl (int msqid, int cmd, struct msqid_ds *buf);
```

Message send

```
#include <sys/msg.h>
```

```
int msgsnd (  
    int msqid, const void *ptr, size_t nbytes, int flag  
);
```

- ❖ Data is placed onto a message queue by calling **msgsnd**
 - The identifier **msqid** specifies the queue on which to send a message
 - The **ptr** argument points to the specific user-defined message data structure **mymsg**
 - **nbytes** specify the size of the data array in **mymsg**

Message send

- The **flag** value is 0 or IPC_NOWAIT
 - If the message queue is full and we specify IPC_NOWAIT, then msgsnd returns with the error **EAGAIN**

❖ Return value

- The value 0, if success
- The value -1, on error

```
int msgsnd (int msqid, const void *ptr,  
            size_t nbytes, int flag);
```

Message receive

```
#include <sys/msg.h>

ssize_t msgrcv (
    int msqid, void *ptr, size_t nbytes, long type,
    int flag
);
```

- ❖ Messages are retrieved from a queue by **msgrcv**
 - Parameters follow the same logic described for **msgsnd**
 - As with **msgsnd**, the **ptr** argument points to the user-defined message structure **mymsg**
 - If the returned message is larger than **nbytes** and the **MSG_NOERROR** bit in **flag** is set, the message is truncated

Message receive

- The type argument lets us specify which message we want

A message queue can implement several FIFOs

- `type=0`: The first message on the queue is returned
- `type>0`: The first message on the queue whose message type equals `type` is returned
- `type<0`: The first message on the queue whose message type is the lowest value less than or equal to the absolute value of `type` is returned

❖ Return value

- Size of data portion of message, if success
- The value, `-1` on error

```
ssize_t msgrcv (int msqid, void *ptr,  
               size_t nbytes, long type, int flag);
```

Example

The Writer

1 Reader + 1 Writer

(W) P_1

(R) P_2

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/ipc.h>
#include <sys/msg.h>

#define L 512

struct msg_buffer {
    long msg_type;
    char msg_text[L];
} message;
```

Example

The Writer

```
int main() {  
    key_t key;  
    int msgid;  
  
    key = ftok ("progfile", 65);  
    msgid = msgget (key, 0666 | IPC_CREAT);  
    message.mesg_type = 1;  
  
    printf ("Read data: ");  
    fgets (message.mesg_text, L, stdin);  
  
    msgsnd (msgid, &message, L*sizeof(char), 0);  
    printf ("Data send: %s\n", message.mesg_text);  
    return 0;  
}
```

Get key

Get the id

Read message
from stdin

Send message

(W) P₁

(R) P₂

Example

The Reader

```
struct mesg_buffer {  
    long mesg_type;  
    char mesg_text[L];  
} message;
```

```
int main() {  
    key_t key;  
    int msgid;  
    key = ftok ("progfile", 65);  
    msgid = msgget (key, 0666 | IPC_CREAT);  
    msgrcv (msgid, &message, L*sizeof(char), 1, 0);  
    printf ("Data received: %s\n", message.mesg_text);  
    msgctl (msgid, IPC_RMID, NULL);  
    return 0;  
}
```

Get key

Get the id

Receive message

Remove the queue
from the system

(W) P₁

(R) P₂

Shared Memory

- ❖ Shared memory allows two or more processes to share a given region of memory
 - It is the fastest form of IPC, because the data does not need to be copied between the client and the server
 - With pipes and message queues, the information has to go through the kernel
 - With shared memory, all processes can access the common memory and changes made by one process can be viewed by another process

Shared Memory

- Shared memory requires synchronization accesses to a given region among multiple processes
 - If the server is placing data into a shared memory region, the client shouldn't try to access the data until the server is done
 - Often, semaphores are used to synchronize shared memory access

Shared Memory

❖ Process logic

- **ftok** is use to generate a unique key to manage the entire process
- **shmget** returns an identifier for the shared memory segment
- **shmat** attach the user to the shared memory segment
- **shmdt** detach the process with with the shared memory segment at the end of the sharing phase
- **shmctl** destroy the shared memory buffer once the process has been detached

Shared memory get

```
#include <sys/shm.h>
```

```
int shmget (key_t key, size_t size, int flag);
```

- ❖ Function **shmget** is used to obtain a shared memory identifier given the **key** of the IPC object
 - The parameter **size** is the size of the shared memory segment in bytes
 - The parameter **flag** set the mode field of the IPC structure
 - See the example for further details

Shared memory get

❖ Return value

- Shared memory ID, if success
- The value -1 , on error

```
int shmget (key_t key, size_t size, int flag);
```

Shared memory control

```
#include <sys/shm.h>
```

```
int shmctl (int shmid, int cmd, struct shmid_ds *buf);
```

- ❖ Function **shmctl** performs various operations on a shared memory
 - The queue is specified by its identifier (msqid)
- ❖ the **cmd** argument specifies a command to be performed on the segment
 - As with function **msgctl**, it is possible to specify
 - PC_STAT
 - IPC_SET
 - IPC_RMID

Shared memory control

- Or, when the process is running in super-user mode
 - SHM_LOCK to lock the shared memory segment in memory
 - SHM_UNLOCK to unlock the shared memory segment

❖ Return value

- 0, if OK
- -1, on error

```
int shmctl (int shmid, int cmd, struct shmid_ds *buf);
```

Shared memory attach

```
#include <sys/shm.h>
```

```
void *shmat (int shmid, const void *addr, int flag);
```

- ❖ Once a shared memory segment has been created, a process attaches it to its address space by calling **shmat**
- ❖ The address in the calling process at which the segment is attached depends on
 - The **addr** argument
 - Whether the SHM_RND bit is specified in **flag** argument

Shared memory attach

Most
common
case

- If **addr** is 0, the segment is attached at the first available address selected by the kernel
- If **addr** is nonzero and SHM_RND is not specified, the segment is attached at the address given by **addr**
- If **addr** is nonzero and SHM_RND is specified, the segment is attached at the address given by $(\text{addr} - (\text{addr} \bmod \text{SHMLBA}))$

❖ Return value

- Pointer to the shared memory segment, if success
- The value -1 , on error

```
void *shmat (int shmid, const void *addr, int flag);
```


Shared memory delete

```
#include <sys/shm.h>

int shmdt (const void *addr);
```

- ❖ When we are done with a shared memory segment, we call **shmdt** to detach it
 - Note that this does not remove the identifier and its associated data structure from the system
 - The identifier must be removed by calling **shmctl** with a command of **IPC_RMID**
- ❖ Return value
 - The value 0, if success
 - The value -1, on error

Example

1 Reader + 1 Writer

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
```

```
#define SHM_SIZE 1024
```

```
int main (int argc, char *argv[]) {
    key_t key;
    int shmid;
    char *data;
```

(W) P₁



(R) P₂

The same process works as

- a reader (no parameter)
- a writer (writing the parameter on the shared memory)

Make it a 1K shared memory segment

Example

Make the
key

Here the file must exist

```
if ((key = ftok ("hello.txt", 5)) == -1) {  
    perror ("ftok");  
    exit (1);  
}  
if ((shmid = shmget (key, SHM_SIZE,  
                    0644 | IPC_CREAT)) == -1) {  
    perror ("shmget");  
    exit (1);  
}  
data = shmat (shmid, NULL, 0);  
if (data == (char *) (-1)) {  
    perror ("shmat");  
    exit (1);  
}
```

Create the segment

Attach the segment to
the local pointer **data**

Example

Modify the segment, based on the command line

```
if (argc == 2) {
    printf ("Writing to segment: \"%s\"\n", argv[1]);
    strncpy (data, argv[1], SHM_SIZE);
}
else
{
    printf("segment contains: \"%s\"\n", data);
}
if (shmdt(data) == -1) {
    perror ("shmdt");
    exit (1);
}
return 0;
}
```

Read the segment

Detach from the segment