

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
{
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE *f;
```

```
for(i=0; i<MAXPAROLA; i++)
    freq[i]=0;
```

```
if(argc != 2)
```

```
{
    printf(stderr, "ERRORE, serve un parametro con il nome del file\n");
    exit(1);
}
```

```
f = fopen(argv[1], "r");
if(f==NULL)
```

```
{
    printf(stderr, "ERRORE, impossibile aprire il file %s\n", argv[1]);
    exit(1);
}
```

```
while( fgets( riga, MAXRIGA, f ) != NULL )
```



High Level Parallel Programming

Programming with the STL

Alessandro Savino

Dipartimento di Automatica e Informatica

Politecnico di Torino

The Standard Library

- ❖ Provides a collection of useful C++ classes and functions
 - Is itself implemented in C++
 - Part of the ISO C++ standard
 - Defines interface, semantics and contracts the implementation has to abide by (e.g. runtime complexity)
 - Implementation is not part of the standard
 - Multiple vendors provide their own implementations
 - Best known
 - **libstdc++** (used by gcc)
 - **libc++** (used by llvm)

The Standard Library

- All features are declared within the `std` namespace
- Functionality is divided into sub-libraries each consisting of multiple headers
- Includes parts of the C standard library
 - For backward compatibility
 - Headers begin with "c" (e.g., `cstring`)
 - Never use them unless you absolutely know what you are doing!

Features Overview

❖ Most important library features

➤ Utilities

- Memory management (new, delete, unique_ptr, shared_ptr)
- Error handling (exceptions, assert)
- Time (clocks, durations, timestamps, etc.)
- Optionals, Variants, Tuples, etc.

➤ Strings

- String class
- String views
- C-style string handling

➤ Containers

- Array, vector, lists, maps, sets

Features Overview

➤ Algorithms

- (Stable) sort, search, max, min, ...

➤ Iterators

➤ Numerics

- Common mathematic functions (sqrt, pow, mod, log, etc.)
- Complex numbers
- Random number generation

➤ I/O

- Input-/Output streams
- File streams
- String streams

Features Overview

➤ Threads

- Thread class
- (shared) mutexes
- Futures

➤ Much more

- Localization
- Regex
- Atomics
- Filesystem support
- ...

Optionals

- ❖ `std::optional` is a class encapsulating a value that might or might not exist
 - Defined in the header `<optional>`
 - Some functions might fail or return without a valid result (e.g. looking up a non-existing file)
 - It's unfavorable to encode such failures with a value of the function domain (e.g., an empty string when file could not be read)
 - It helps to express such results
 - At any point in time, an optional either has a value, or it doesn't
 - If the computation succeeded, it returns an optional containing a value

Optionals

- If it failed, it returns an optional without a value
- The template parameter T denotes, of which type the optional may contain a value
 - For example, `optional<int>` might contain an int
- Guarantees to not dynamically allocate any memory when being assigned a value
- Is an object, despite supporting the dereference operators `*` and `->`
- Internally implemented as an object with a member of type T and a boolean

Optionals: Creations

```
std::optional<std::string> might_fail(int arg) {  
    if (arg == 0) {  
        return std::optional<std::string>("zero");  
    } else if (arg == 1) {  
        return "one"; // equivalent to the case above  
    } else if (arg < 7) {  
        // std::make_optional takes constructor arguments of type T  
        return std::make_optional<std::string>("less than 7");  
    } else {  
        return std::nullopt; // alternatively: return {}  
    }  
}
```

- ❖ Optionals are created through its constructor or with `std::make_optional`

Optionals: Creations

```
might_fail(3).value(); // "less than 7"

might_fail(8).value();
// throws std::bad_optional_access

*might_fail(3); // "less than 7"

might_fail(6)->size(); // 11

might_fail(7)->empty(); // undefined behavior
```

- ❖ The value of an optional can be read with `value()` (throws exception when empty) or dereferenced with `*` or `->` (undefined behavior when empty)

Optionals: Checking and Accessing

```
might_fail(3).has_value(); // true
might_fail(8).has_value(); // false

// Or even simpler:
std::optional<std::string> opt5 = might_fail(5)
if (opt5) { //contextual conversion to bool
    opt5->size(); // 11
}
```

- ❖ There are multiple ways to check whether an optional has a value

Pairs

- ❖ `std::pair<T, U>` is a template class that stores exactly one object of type `T` and one of type `U`
 - Defined in the header `<utility>`
 - Constructor takes object of `T` and `U`
 - Pairs can also be constructed with `std::make_pair()`
 - Objects can be accessed with `first` and `second`
 - Can be compared for equality and inequality
 - Can be compared lexicographically with `<`, `<=`, `>`, and `>=`

```
std::pair<int, double> p1(123, 4.56);  
p1.first; // == 123  
p1.second; // == 4.56
```

```
auto p2 = std::make_pair(456, 1.23);  
// p2 has type std::pair<double, int>  
p1 < p2; // true
```

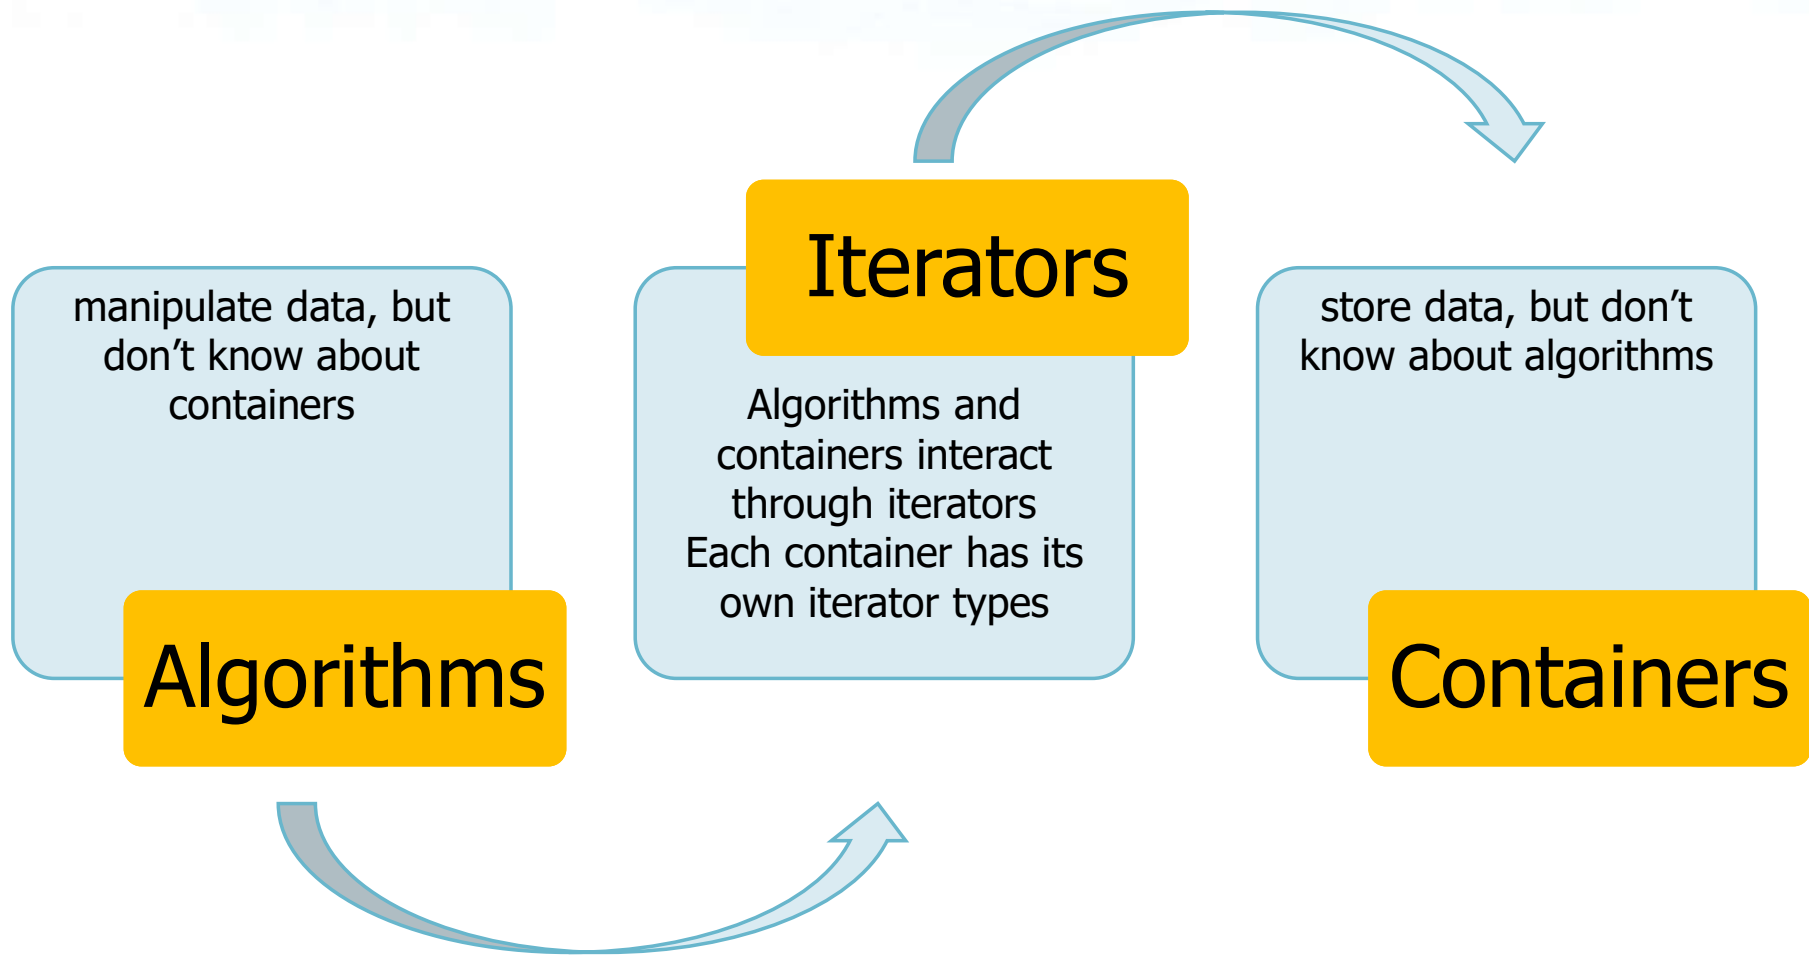
Tuple

- ❖ `std::tuple` is a template class with `n` type template parameters that stores exactly one object of each of the `n` types.
 - Defined in the header `<tuple>`
 - Constructor takes all objects
 - Tuples can also be constructed with `std::make_tuple()`
 - The `i`th object can be accessed with `std::get<i>()`
 - Just like pairs, tuples define all relational comparison operators

```
std::pair<int, double, char> t1(123, 4.56, 'x');  
std::get<1>(t1); // == 4.56
```

```
auto t2 = std::make_tuple(456, 1.23, 'y');  
// t2 has type std::tuple<int, double, char>  
t1 < t2; // true
```

The STL Basic Model



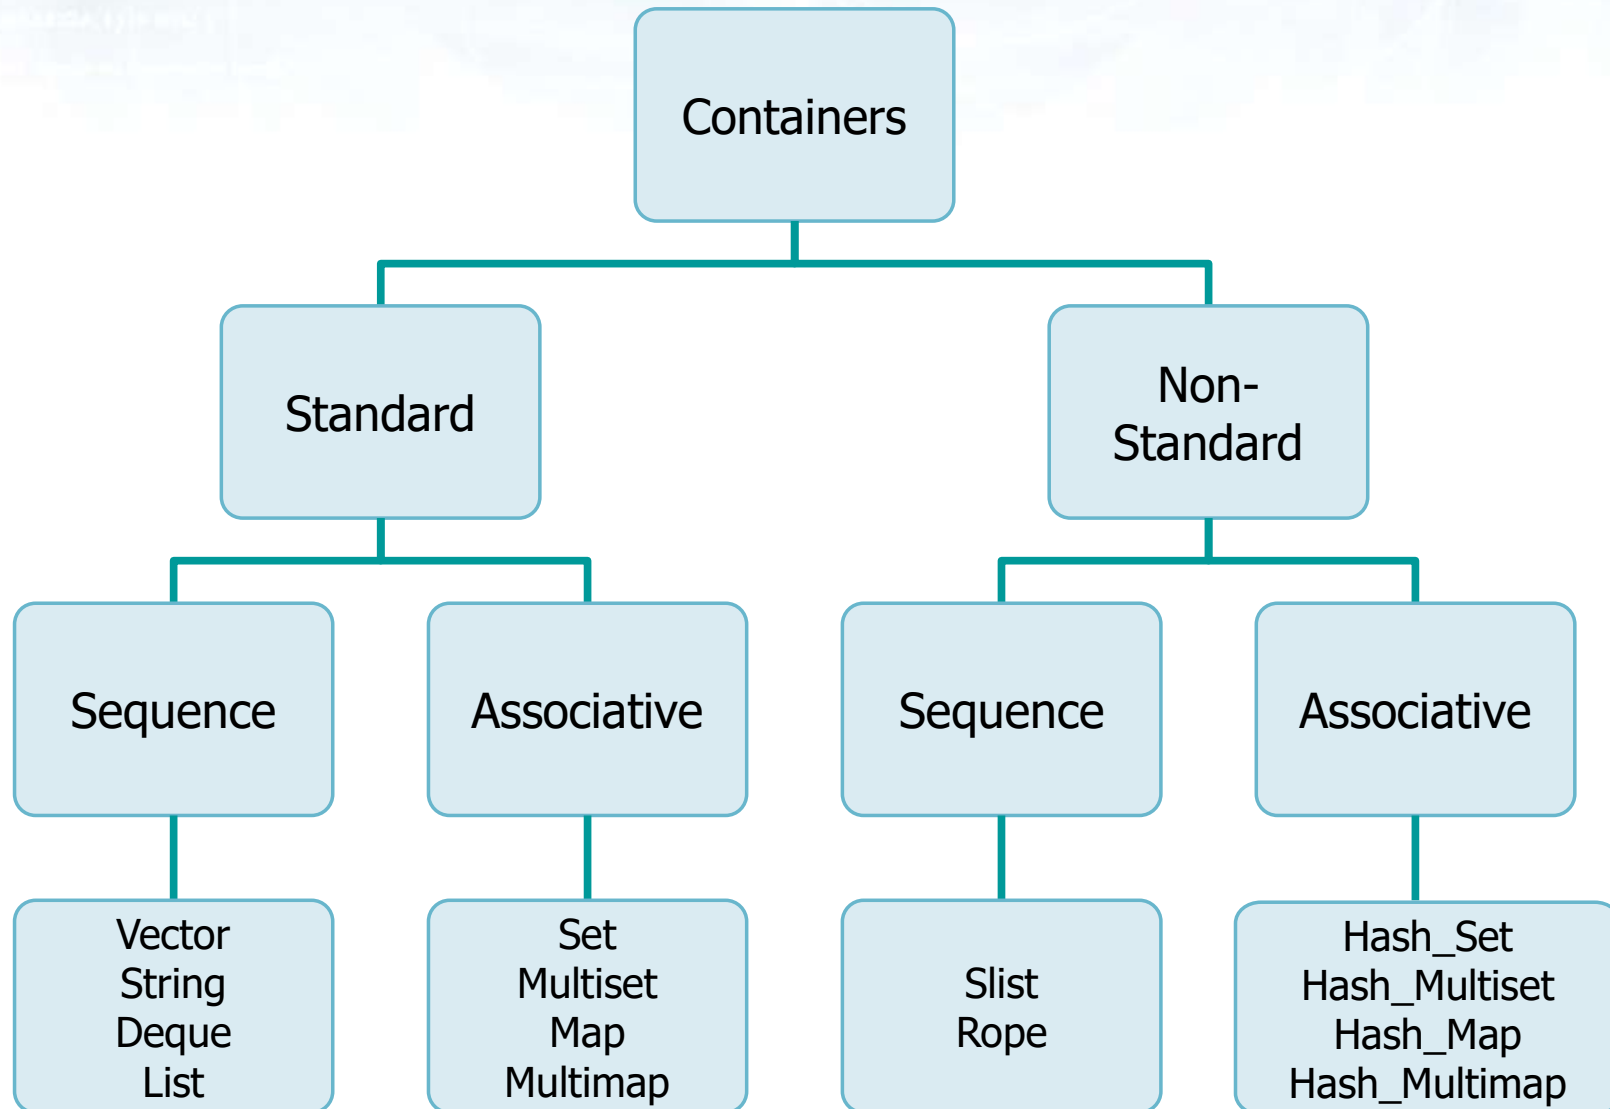
Containers

- ❖ A container is an object that stores a collection of other objects
 - Manage the storage space for their elements
 - Generic
 - The type(s) of elements stored are template parameter(s)
 - Provide member functions for accessing elements directly, or through iterators
 - (Most) member functions shared between containers

Containers

- Make guarantees about the complexity of their operations
 - Sequence containers (e.g. `std::array`, `std::vector`, `std::list`): Optimized for sequential access
 - Associative containers (e.g. `std::set`, `std::map`): Sorted, optimized for search ($O(\log n)$)
 - Unordered associative containers (e.g. `std::unordered_set`, `std::unordered_map`): Hashed, optimized for search
 - Amortized: $O(1)$
 - Worst case: $O(n)$
- ❖ Use containers whenever possible
 - When in doubt, use `std::vector`!

Types of Containers



Containers: `std::vector`

- ❖ Vectors are arrays that can dynamically grow
 - Defined in the header `<vector>`
 - Elements are still stored contiguously
 - Elements can be inserted and removed at any position
 - Pre-allocates memory for a certain amount of elements
 - Allocates new, larger chunk of memory and moves elements when memory is exhausted
 - Memory for a given amount of elements can be reserved with `reserve`

Containers: `std::vector`

- Time complexity
 - Random access
 - $O(1)$
 - Insertion and removal at the end
 - Typically: $O(1)$
 - Worst case: $O(n)$ due to possible reallocation
 - Insertion and removal at any other position
 - $O(n)$
- Access to the underlying C-style data array with **data** member function

Containers: `std::vector`

- ❖ Vectors are constructed just like arrays
- ❖ Access elements via C-style array notation, via `at()`, or through a raw pointer
- ❖ Update elements via C-style array notation, via `at()`, or through a raw pointer

➤ Note

- It is not possible to insert new elements this way
- You can only update existing ones

```
std::vector<int> fib = {1,1,2,3};

fib.at(0) // == 1;
fib[1] // == 1;
int* fib_ptr = fib.data();
fib_ptr[2] // == 3;

fib[3] = 43;
fib.at(2) = 42;
fib.data()[1] = 41; // fib is now 1, 41, 42, 43
```

Containers: `std::unordered_map`

- ❖ Maps are associative containers consisting of key-value pairs
 - Defined in the header `<unordered_map>`
 - Keys are required to be unique
 - At least two template parameters
 - Key and T (type of the values)
 - Is internally a hash table
 - Amortized **$O(1)$** complexity for random access, search, insertion, and removal

Containers: `std::unordered_map`

- No way to access keys or values in order (use `std::map` for that!)
- Accepts custom hash- and comparison functions through third and fourth template parameter
- ❖ Use `std::unordered_map` if you need a hash table, but don't need ordering

Containers: `std::unordered_map`

- ❖ Maps can be constructed pairwise
- ❖ Lookup the value to a key with C-style array notation, or with `at()`
- ❖ A pair can also be searched for with `find`
 - To check if a key exists, use `count`

```
std::unordered_map<std::string, double> name_to_grade
{{"maier", 1.3}, {"huber", 2.7}, {"schmidt", 5.0}};

name_to_grade["huber"]; // == 2.7
name_to_grade.at("schmidt"); // == 5.0

auto search = name_to_grade.find("schmidt");
if (search != name_to_grade.end()) {
    // Returns an iterator pointing to a pair!
    search->first; // == "schmidt"
    search->second; // == 5.0
}

name_to_grade.count("schmidt"); // == 1
name_to_grade.count("blafasel"); // == 0
```

Containers: `std::map`

- ❖ In contrast to unordered maps, the keys of `std::map` are sorted
 - Defined in the header `<map>`
 - Interface largely the same to `std::unordered_map`
 - Optionally accepts a custom comparison function as template parameter
 - Is internally a tree (usually AVL- or R/B-Tree)
 - $O(\log n)$ complexity for random access, search, insertion, and removal

Containers: `std::map`

- ❖ `std::map` also allows to search for ranges
 - `upper_bound(key)` returns an iterator pointing to the first **greater** element than key
 - `lower_bound(key)` returns an iterator pointing to the first element **not lower** than key

```
std::map<int, int> x_to_y = {{1, 1}, {3, 9}, {7, 49}};  
  
// gt3 points to {7, 49}  
auto gt3 = x_to_y.upper_bound(3);  
  
// geq3 points to {3, 9}  
auto geq3 = x_to_y.lower_bound(3);
```

Containers: Thread Safety

- ❖ Containers give some thread safety guarantees:
 - Two different containers: All member functions can be called concurrently by different threads (i.e. different containers don't share state)
 - The same container: All const member functions can be called concurrently. `at()`, `[]` (except in associative containers), `data()`, `front()/back()`, `begin()/end()`, `find()` also count as const
 - Iterator operations that only read (e.g. incrementing or dereferencing an iterator) can be run concurrently with reads of other iterators and const member functions
 - Different elements of the same container can be modified concurrently

Containers: Thread Safety

❖ Be careful

- As long as the standard does not explicitly require a member function to be sequential, the standard library implementation is allowed to parallelize it internally (see e.g. `std::transform` vs. `std::for_each`)

❖ Rule of thumb

- Simultaneous reads on the same container are always okay
- Simultaneous read/writes on different containers are also okay
- Everything else requires synchronization

Iterators

❖ Iterators are objects that can be thought of as pointer abstractions

➤ Problem

- Different element access methods for each container

➤ Therefore

- Container types not easily exchangeable in code

➤ Solution

- Iterators abstract over element access and provide pointer-like interface

Iterators

- Allow for easy exchange of underlying container type
- The standard library defines multiple iterator types as containers have varying capabilities (random access, traversable in both directions, etc.)
- ❖ **Be careful**
 - When writing to a container, all existing iterators are invalidated and can no longer be used (some exceptions apply)

Iterators: An Example

```
std::vector<std::string> vec = {"one", "two", "three", "four"};

std::vector<std::string>::iterator it = vec.begin();
auto end = vec.end();

std::cout << *it; // prints "one"
std::cout << it->size(); // prints 3

std::cout << *end; // undefined behavior

++it; // Prefer to use pre-increment
std::cout << *it; // prints "two"

// prints "three,four,"
while (it != end) {
    std::cout << *it << ",";
    it++;
}

// available also in the range expression (C++11 and +)
for (auto elem : vec) {
    std::cout << elem << ","; // prints "one,two,three,four,"
}
```

Iterators: An Example

```
std::vector<std::string> vec = {"one", "two", "three", "four"};

std::vector<std::string>::iterator it = vec.begin();
auto end = vec.end();

std::cout << "vec: ";
std::cout << " ";
std::cout << " ";
++it; // Pre
std::cout << " ";

// prints "t
while (it != end) {
    std::cout << *it <
    it++;
}

// available also in the range expression (C++11 and +)
for (auto elem : vec) {
    std::cout << elem << ","; // prints "one,two,three,four,"
}
```

- Such a loop requires the range expression (here: vec) to have a begin() and end() member.
- vec.begin() is assigned to an internal iterator which is dereferenced, assigned to the range declaration (here: auto elem), and then incremented until it equals vec.end().

Iterators: An Advanced Example

```
std::vector<std::string> vec = {"one", "two", "three", "four"};

for (it = vec.begin(); it != vec.end(); ++it) {
    // (*it).size
    if (it->size == 3) {
        it = vec.insert(it, "foo");
        // it now points to the newly inserted element
        ++it;
    }
}

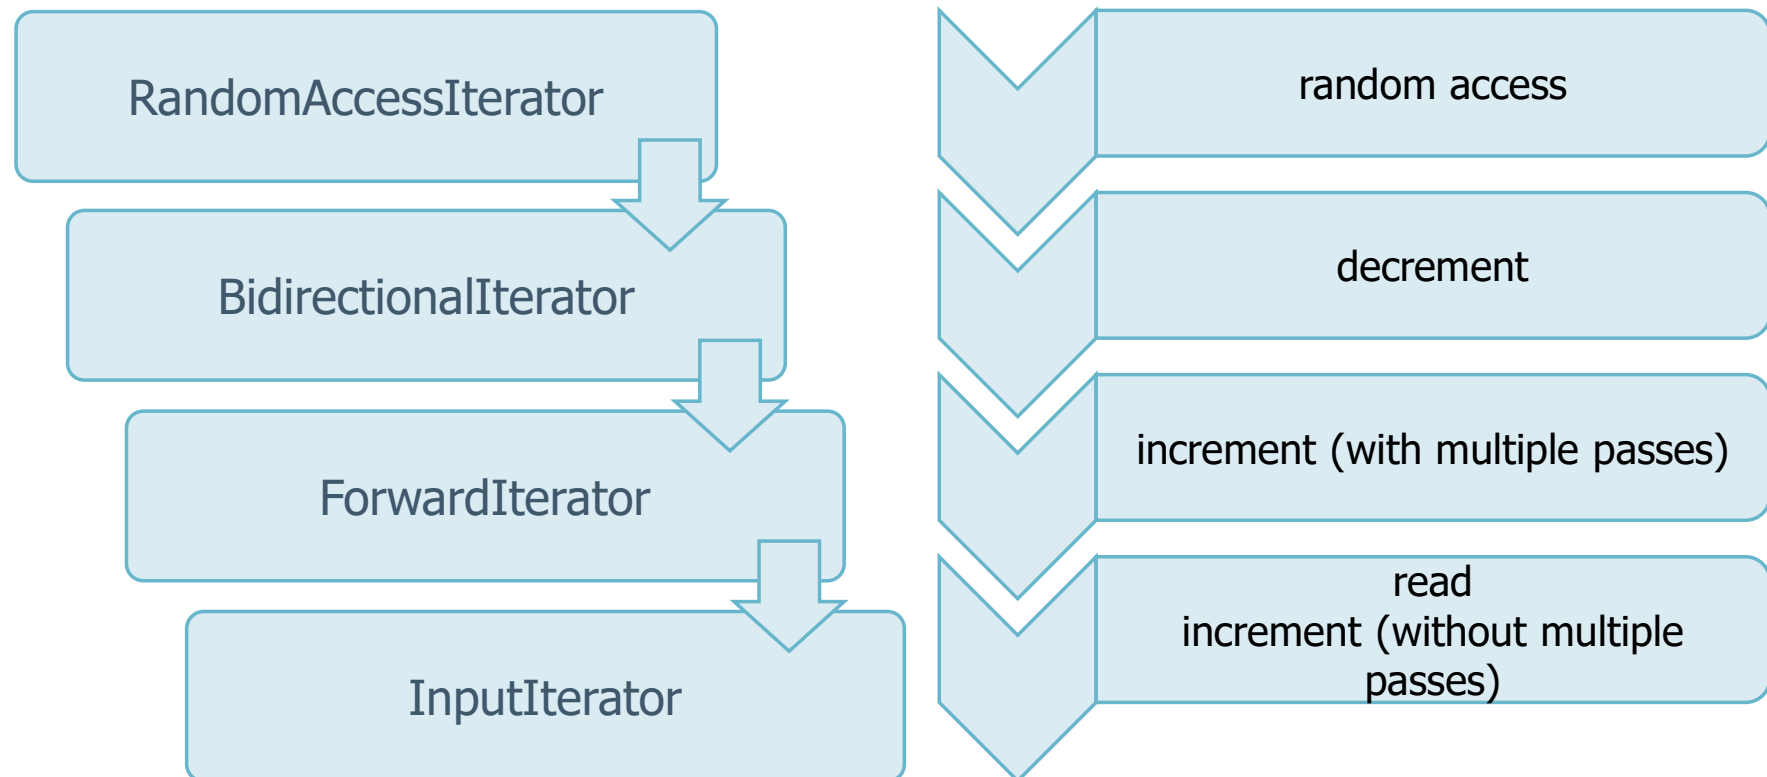
//vec == {"foo", "one", "foo", "two", "three", "four"}

for (it = vec.begin(); it != vec.end(); ++it) {
    if (it->size == 3) {
        it = vec.erase(it);
        // erase returns a new, valid iterator
        // pointing at the next element
    }
}

//vec == {"three", "four"}
```


Iterators Hierarchy

- ❖ Iterators are instantiation of classes
 - It exist a hierarchy of Iterators
 - They differ by the kind of operations allowed



Iterators Hierarchy

- ❖ Iterators are instantiation of classes
 - It exist a hierarchy of Iterators
 - They differ by the kind of operations allowed

OutputIterator



write
increment (without
multiple passes)

Iterators: Input and Output

- ❖ Input- and OutputIterator are the most basic iterators
- ❖ They have the following features
 - Equality comparison
 - Checks if two iterators point to the same position
 - Dereferencable with the * and -> operators
 - Incrementable, to point at the next element in sequence
 - A dereferenced **InputIterator** can only be read
 - A dereferenced **OutputIterator** can only be written to

Iterators: Input and Output

- ❖ As the most restrictive iterators, they have a few limitations
 - Single-pass only
 - They cannot be decremented
 - Only allow equality comparison, $<$, $>=$, etc. not supported
 - Can only be incremented by one (i.e., $it + 2$ does not work)
- ❖ Used in single-pass algorithms such as `find()` (InputIterator) or `copy()` (Copying from an InputIterator to an OutputIterator)

Iterators: Forward and Bidirectional

- ❖ ForwardIterator combines InputIterator and OutputIterator
 - All the features and restrictions shared between input- and output iterator apply
 - Dereferenced iterator can be read and written to
- ❖ BidirectionalIterator generalizes ForwardIterator
 - Additionally allows decrementing (walking backwards)
 - Therefore supports multi-pass algorithms traversing the container multiple times
 - All other restrictions of ForwardIterator still apply

Iterators: Random

- ❖ RandomAccessIterator generalizes BidirectionalIterator
 - Additionally allows random access with operator[]
 - Supports relational operators, such as < or >=
 - Can be incremented or decremented by any amount (i.e. it + 2 does work)

Function Objects

- ❖ Regular functions are not objects in C++
 - Cannot be passed as parameters (unless a function pointer is allowed)
 - Cannot have state
- ❖ C++ additionally defines the FunctionObject named requirement
- ❖ For a type T to be a FunctionObject
 - T has to be an object
 - operator()(args) has to be defined for T for a suitable argument list args which can be empty
 - Often referred to as **functors**

Function Objects

- ❖ There are several valid function objects defined in C++
 - Pointers to functions
 - Lambda expressions
 - Stateful function objects in form of classes
- ❖ Functions and function references are not function objects
 - Can still be used in the same way due to implicit function-to-pointer conversion

Function Pointers

- ❖ While functions are not objects, they do have an address
 - Location in memory where the actual assembly code resides
 - Allows declaration of function pointers
- ❖ Function pointers to non-member functions
 - Declaration
 - return-type (*identifier)(args)
 - Allows passing functions as parameters
 - E.g., passing a custom compare function to **std::sort**
 - E.g., passing a callback to a method
- ❖ Can be invoked in the same way as a function

Function Pointers

```
int callFunc(int (*func)(int, int), int arg1, int arg2) {
    return (*func)(arg1, arg2);
}

//-----
double callFunc(double (*func)(double), double argument) {
    return func(argument); // Automatically dereferenced
}

//-----
int add(int arg1, int arg2) { return arg1 + arg2; }
double add4(double argument) { return argument + 4; }

//-----
int main() {
    auto i = callFunc(add, 2, 4); // i = 6
    auto j = callFunc(&add4, 4); // j = 8, "&" can be omitted
}
```

Lambda Expressions

- ❖ Function pointers can be unwieldy
 - Function pointers cannot easily capture environment
 - Must pass all variables that affect function by parameter
 - Cannot have “local” functions within other functions
- ❖ C++ defines lambda expressions as a more flexible alternative
 - Lambda expressions construct a closure
 - Closures store a function together with an environment
 - Lambda expressions can capture variables from the scope where they are defined

Lambda Expressions

❖ Lambda expression syntax

```
[ captures ] ( params ) -> ret { body }
```

- **captures** specifies the parts of the environment that should be stored
- **params** is a comma-separated list of function parameters
- **ret** specifies the return type and can be omitted, in which case the return type is deduced from return statements inside the **body**

Lambda Expressions

- ❖ The list of captures can be empty
 - Results in stateless lambda expression
 - Stateless lambda expressions are implicitly convertible to function pointers
- ❖ Lambda expressions have unique unnamed class type
 - Must use auto when assigning lambda expressions to variables
 - Declaration of a lambda variable (e.g. as member) is not possible

Lambda Expression

```
int callFunc(int (*func)(int, int), int arg1, int arg2) {  
    return func(arg1, arg2);  
}
```

```
//-----  
int main() {  
  
    auto lambda = [](int arg1, int arg2) {  
        return arg1 + arg2;  
    };  
  
    int i = callFunc(lambda, 2, 4); // i = 6  
  
    int j = lambda(5, 6); // j = 11  
}
```

Lambda Expression: Uniqueness constraint

```
// ERROR: Compilation will fail due to ambiguous return type
auto getFunction(bool first) {
    if (first) {
        return []() {
            return 42;
        };
    } else {
        return []() {
            return 42;
        };
    }
}
```

All lambda expressions have **unique** types

Lambda Expressions: Captures

- ❖ Lambda captures specify what constitutes the state of a lambda expression
 - Can refer to *automatic variables* in the surrounding scopes (up to the enclosing function)
 - Can refer to the “this” pointer in the surrounding scope (if present)
- ❖ Captures can either capture *by-copy* or *by-reference*
 - Capture by-copy creates a copy of the captured variable in the lambda state
 - Capture by-reference creates a reference to the captured variable in the lambda state
 - Captures can be used in the lambda expression body like regular variables or references

Lambda Expressions: Captures

- ❖ Lambda captures are provided as a comma-separated list of captures
 - By-copy
 - *identifier* or *identifier initializer*
 - By-reference
 - *&identifier* or *&identifier initializer*
 - *identifier* must refer to automatic variables in the surrounding scopes
 - *identifier* can be used as an identifier in the lambda body
 - Each variable may be captured only once

Lambda Expressions: Captures

❖ First capture can optionally be a capture-default

➤ By-copy

- =

➤ By-reference

- &

- Allows any variable in the surrounding scopes to be used in the lambda body
- Specifies the capture type for all variables without explicit captures
- If present, only diverging capture types can be specified afterwards

Lambda Expression: Captures

```
int main() {  
    int i = 0;  
    int j = 42;  
  
    auto lambda1 = [i](){};           // i by-copy  
    auto lambda2 = [&i](){};          // i by-reference  
  
    auto lambda3 = [&j, i](){};       // j by-reference, i by-copy  
    auto lambda4 = [=, &i](){};      // j by-copy, i by-reference  
  
    // ERROR: non-diverging capture types  
    auto lambda5 = [&, &i](){};  
    // ERROR: non-diverging capture types  
    auto lambda6 = [=, i](){};  
}
```

Lambda Expression: Captures

```
int main() {  
    int i = 42;  
    auto lambda1 = [i]() { return i + 42; };  
    auto lambda2 = [&i]() { return i + 42; };  
  
    i = 0;  
    int a = lambda1(); // a = 84  
    int b = lambda2(); // b = 42  
}
```

Be careful, the capture is done at definition, thus if you capture by-copy the value is persistent, otherwise (by-reference) the reference is persistent

Lambda Expression: Captures

```
#include <memory>

int main() {
    auto ptr = std::make_unique<int>(4);
    auto f2 = [inner = ptr.get()]() { return *inner;};

    int a = f2(); //4
    ptr.reset();
    int b = f2(); // undefined behavior
}
```

Lifetime issues due to capturing...

Stateful Function Objects

- ❖ Situation so far
 - Functions are generally stateless
 - State must be kept in surrounding object, e.g., class instances
 - Lambda expressions allow limited state-keeping
- ❖ Function objects can be implemented in a regular class
 - Allows the function object to keep arbitrary state
 - Difference to lambda expressions
 - State can be changed during lifetime

Stateful Function Objects: `std::function`

- ❖ `std::function` is a general-purpose wrapper for all callable targets
 - Defined in the `<functional>` header
 - Able to store, copy and invoke the wrapped target
 - Potentially incurs dynamic memory allocations
 - Often adds unnecessary overhead
 - **Should be avoided where possible**

Stateful Function Objects: `std::function`

```
#include <functional>

//-----
std::function<int()> getFunction(bool first){
    int a = 14;
    if (first)
        return [=]() { return a; };
    else
        return [=]() { return 2 * a; };
}

//-----
int main() {
    return getFunction(false)() + getFunction(true)(); // 42
}
```


The Algorithms Library

- ❖ The algorithms library is part of the C++ standard library
 - Defines operations on ranges of elements [first, last)
 - Bundles functions for sorting, searching, manipulating, etc.
 - Ranges can be specified using pointers or any appropriate iterator type
 - Spread in 4 headers
 - `<algorithm>`
 - `<numeric>`
 - `<memory>`
 - `<cstdlib>`
- ❖ We will focus on `<algorithm>` as it bundles the most relevant parts

std::sort

- ❖ Sorts all elements in a range [first, last) in ascending order
 - `void sort(RandomIt first, RandomIt last);`
 - Iterators must be RandomAccessIterators
 - Elements must be swappable (`std::swap` or user-defined swap)
 - Elements must be move-assignable and move-constructible
 - Does not guarantee order of equal elements
 - Needs $O(n * \log(N))$ comparisons

std::sort

```
// using default comparison
// operator
#include <algorithm>
#include <vector>

//-----
int main() {
    std::vector<unsigned> v =
        {3, 4, 1, 2};
    std::sort(v.begin(), v.end());
    // 1, 2, 3, 4
}
```

```
// using custom comparison
#include <algorithm>
#include <vector>

//-----
int main() {
    std::vector<unsigned> v =
        {3, 4, 1, 2};
    std::sort(v.begin(),
        v.end(),
        [](unsigned lhs,
            unsigned rhs) {
                return lhs > rhs;
            }); // 4, 3, 2, 1
}
```

std::sort

```
// using default comparison
// operator
#include <algorithm>
#include <vector>
```

```
//-
int
```

- Sorting algorithms can be modified through custom comparison functions
 - Supplied as function objects (Compare named requirement)
- Must establish a strict weak ordering on the elements

- Syntax:

```
bool cmp(const Type1 &a, const Type2
        &b);
```

- Return true if and only if a and b must be swapped

```
// using custom comparison
#include <algorithm>
#include <vector>

//-----
int main() {
    std::vector<unsigned> v =
        {3, 4, 1, 2};
    std::sort(v.begin(),
              v.end(),
              [](const unsigned &lhs,
                const unsigned &rhs) {
                  return lhs > rhs;
              }); // 4, 3, 2, 1
}
```

Further Sorting

- ❖ Sometimes `std::sort` may not be the optimal choice
 - Does not necessarily keep order of equal-ranked elements
 - Sorts the entire range (unnecessary e.g. for top-k queries)
- ❖ Keep the order of equal-ranked elements
 - `std::stable_sort`
- ❖ Partially sort a range
 - `std::partial_sort`
- ❖ Check if a range is sorted
 - `std::is_sorted`
 - `std::is_sorted_until`

Searching

- ❖ The algorithms library offers a variety of searching operations
 - Different set of operations for sorted and unsorted ranges
 - Searching on sorted ranges is faster in general
 - Sorting will pay off for repeated lookups
- ❖ Arguments against sorting
 - Externally prescribed order that may not be modified
 - Frequent updates or insertions
- ❖ General semantics
 - Search operations return iterators pointing to the result
 - Unsuccessful operations are usually indicated by returning the end iterator of a range [first, last)

Searching

- ❖ Find the first element satisfying some criteria
 - `std::find`
 - `std::find_if`
 - `std::find_if_not`
- ❖ Search for a range of elements in another range of elements
 - `std::search`
- ❖ Count matching elements
 - `std::count`
 - `std::count_if`
- ❖ Many more useful operations (see reference documentation)

std::find, std::find_if

```
#include <algorithm>
#include <vector>

//-----

int main() {
    std::vector<int> v = {2, 6, 1, 7, 3, 7};
    auto res1 = std::find(vec.begin(), vec.end(), 7);
    int a = std::distance(vec.begin(), res1);
    // a = 3 is the index distance between Iterator begin() and res1

    auto res2 = std::find(vec.begin(), vec.end(), 9);
    if(res2 == vec.end())
        std::cout << "Not found!";

    auto res1 = std::find_if(vec.begin(),
                            vec.end(),
                            [](int val) { return (val % 2) == 1; }
                            );
    int a = std::distance(vec.begin(), res1); // 2
}
```


Searching: sorted ranges

- ❖ On sorted ranges, binary search operations are offered
 - Complexity $O(\log(N))$ when range is given as `RandomAccessIterator`
 - Can employ custom comparison function (see above)
 - When called with `ForwardIterators` complexity is linear in number of iterator increments
- ❖ Search for one occurrence of a certain element
 - `std::binary_search`
- ❖ Search for range boundaries
 - `std::lower_bound`
 - `std::upper_bound`
- ❖ Search for all occurrences of a certain element
 - `std::equal_range`
- ❖ Certain Containers already implement (some of) them!

Permutations

- ❖ The algorithms library offers operations to permute a given range
 - Can iterate over permutations in lexicographical order
 - Requires at least BidirectionalIterators
 - Values must be swappable
 - Order is determined using operator< by default
 - A custom comparison function can be supplied (see above)
- ❖ Initialize a dense range of elements
 - `std::iota`
- ❖ Iterate over permutations in lexicographical order
 - `std::next_permutation`
 - `std::prev_permutation`

Additional Functionalities

- ❖ The algorithms library offers many more operations
 - `std::min` & `std::max` over a range instead of two elements
 - `std::merge` & `std::in_place_merge` for merging of sorted ranges
 - Multiple set operations (intersection, union, difference, ...)
 - Heap functionality
 - Sampling of elements using `std::sample`
 - Swapping elements using `std::swap`
 - Range modifications
 - `std::copy` To copy elements to new location
 - `std::rotate` To rotate range
 - `std::shuffle` To randomly reorder elements
- ❖ For even more operations see the reference documentation