

```
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
```

```
#define MAXPAROLA 30
#define MAXRIGA 80
```

```
int main(int argc, char *argv[])
{
    int freq[MAXPAROLA]; /* vettore di contatori
delle frequenze delle lunghezze delle parole */
    char riga[MAXRIGA];
    int i, inizio, lunghezza;
    FILE *f;
```

```
for(i=0; i<MAXPAROLA; i++)
    freq[i]=0;
```

```
if(argc != 2)
```

```
{
    fprintf(stderr, "ERRORE, serve un parametro con il nome del file\n");
    exit(1);
}
```

```
f = fopen(argv[1], "r");
if(f==NULL)
```

```
{
    fprintf(stderr, "ERRORE, impossibile aprire il file %s\n", argv[1]);
    exit(1);
}
```

```
while( fgets( riga, MAXRIGA, f ) != NULL )
```



High Level Parallel Programming

Multi-Threading in C++

Alessandro Savino

Dipartimento di Automatica e Informatica

Politecnico di Torino

Introduction

- ❖ C++ can run multiple threads simultaneously
 - Multiple threads use the same memory
 - They may read from the same memory location
 - All other accesses (i.e. read-write, write-read, write-write) are called conflicts
 - Conflicting operations are only allowed when threads are synchronized
 - This can be done with mutexes or atomic operations
 - Unsynchronized accesses (also called data races), deadlocks, and other potential issues when using threads are undefined behavior

Introduction

- ❖ The header `<thread>` defines the class `std::thread` that can be used to start new threads
 - Using this class is the best way to use threads platform-independently
 - Using it may require additional compiler flags
 - `-pthread` for gcc and clang

```
set(CMAKE_CXX_FLAGS "${CMAKE_CXX_FLAGS} -std=c++14 -pthread")
```

Example

```
void foo(int a, int b);
```

Starts a thread that
calls foo(123, 456)

```
std::thread t1(foo, 123, 456);
```

Also works
with lambdas

```
std::thread t2([] { foo(123, 456); });
```

```
std::thread t3;
```

Creates an object that
does not refer to a thread

Join Threads

- ❖ The member function **join** can be used to wait for a thread to finish
 - Function **join** must be called exactly once for each thread
 - When the destructor of a **std::thread** is called, the program is terminated if it has an associated thread that was not joined

Example

```
std::vector<std::thread> threadPool;  
for (int i = 1; i <= 9; ++i) {  
    threadPool.emplace_back([i] { safe_print(i); });  
}  
  
for (auto& t : threadPool) {  
    t.join();  
}
```

Digits 1 to 9 are
printed (unordered)

Moving Threads

- ❖ **std::threads** are not copyable, but movable, so they can be used in containers
 - Moving an **std::thread** transfers all resources associated with the running thread
 - Only the moved-to thread can be joined

Example

```
std::thread t1([] { std::cout << "Hi\n"; });  
std::thread t2 = std::move(t1);  
t2.join();
```

The thread originally
started in t1 is joined

t1 is now empty

Other useful functions

- ❖ The thread library also contains other useful functions that are closely related to starting and stopping threads
 - `std::this_thread::sleep_for`
 - Stop the current thread for a given amount of time
 - `std::this_thread::sleep_until`
 - Stop the current thread until a given point in time
 - `std::this_thread::yield`
 - Let the operating system schedule another thread
 - `std::this_thread::get_id`
 - Get the (operating-system-specific) id of the current thread

Exercise

- ❖ Write a simple program that evaluates all prime numbers in a given range
 - Range is passed to the program as a parameter
 - Set the number of parallel threads using
 - `std::thread::hardware_concurrency`
 - It reports the actual max number of threads based on your architecture: get along and find out performances

Mutual Exclusion

- ❖ When working with threads, mutual exclusion is a central concept to synchronize threads
- ❖ The standard library defines several useful classes for this in `<mutex>` and `<shared_mutex>`
 - `std::mutex`
 - Mutual exclusion
 - `std::recursive_mutex`
 - Recursive mutual exclusion
 - `std::shared_mutex`
 - Mutual exclusion with shared locks

Mutual Exclusion

- `std::unique_lock`
 - RAII wrapper for `std::mutex`
- `std::shared_lock`
 - RAII wrapper for `std::shared_mutex`
- ❖ **Note**
 - Mutexes are usually inefficient as they are used very coarse-grained and sometimes require communication with the operating system

Mutexes

- ❖ A mutex is the most basic synchronization primitive which can be locked and unlocked by exactly one thread at a time
 - **std::mutex** has the member functions
 - **lock** and **unlock** that lock and unlock the mutex
 - **try_lock** that tries to lock the mutex and returns true if it was successful
 - All three functions may be called simultaneously by different threads
 - For each call to **lock** the same thread must call **unlock** exactly once

Example

```
std::mutex printMutex;

void safe_print(int i) {
    printMutex.lock();
    std::cout << i;
    printMutex.unlock();
}

int main() {
    thread t1(safe_print), t2(safe_print);

    t1.join();
    t2.join();
}
```

Recursive Mutexes

- ❖ Recursive mutexes are regular mutexes that additionally allow a thread that currently holds the mutex to lock it again
 - Implemented in the class **std::recursive_mutex**
 - Has the same member functions as **std::mutex**
 - **unlock** must still be called once for each **lock**
 - Useful for functions that call each other and use the same mutex

Example

```
std::recursive_mutex m;

void foo() {
    m.lock();
    std::cout << "foo\n";
    m.unlock();
}

void bar() {
    m.lock();
    std::cout << "bar\n";
    foo(); // This will not deadlock
    m.unlock();
}
```


Shared Mutexes

- ❖ A shared mutex is a mutex that differentiates between shared and exclusive locks
 - Implemented in the class **std::shared_mutex**
 - A shared mutex can either be locked exclusively by one thread or have multiple shared locks
 - The member functions **lock** and **unlock** are exclusive
 - The member functions **lock_shared** and **unlock_shared** are shared

Since C++17 !

Shared Mutexes

- The member functions **try_lock** and **try_lock_shared** try to get an exclusive or shared lock and return true on success
- ❖ Shared mutexes are mostly used to implement read/write-locks
 - Readers use shared locks, writers use exclusive locks

Example

```
int value = 0;
std::shared_mutex m;

std::vector<std::thread> threadPool;

// Add readers
for (int i = 0; i < 5; ++i)
    threadPool.emplace_back([&] {
        m.lock_shared();
        safe_print(value);
        m.unlock_shared();
    })

// Add writers
for (int i = 0; i < 5; ++i)
    threadPool.emplace_back([&] {
        m.lock(); ++value; m.unlock();
    })
```

Join and other coordination mechanisms still needed

Try this with POSIX mutex !

Usage Rules

- ❖ Mutexes have several requirements on how they must be used
 - For each call to **lock**, **unlock** must be called exactly once
 - **unlock** must only be called by the thread that called **lock**
 - The above also holds for **unlock_shared** and **lock_shared**
 - A thread A should not wait for a mutex from thread B to be unlocked if B needs to lock a mutex that A is currently holding (i.e. avoid deadlocks)

Usage Rules

- ❖ Note the following when using mutexes to make data structures thread-safe
 - The member functions **lock** and **unlock** are non-const
 - If const member functions of the data structure should also use the mutex, it should be mutable
 - If a member function that locks the mutex calls other member functions, this can lead to deadlocks
 - **recursive_mutex** can be used to avoid this

Mutex RAII Wrappers

- ❖ Mutexes can be thought of resources that must be acquired and freed with **lock** and **unlock**
 - The RAII pattern should be used
 - `std::unique_lock` is an RAII wrapper for Mutexes that calls **lock** in its constructor and **unlock** in its destructor
 - **`std::unique_lock`** is movable to “transfer ownership” of the locked mutex
 - It also has the member functions **lock** and **unlock** to manually control the mutex

Example

```
std::mutex m;  
  
int i = 0;  
  
std::thread t([&] {  
    std::unique_lock l{m};  
    ++i;  
    ...  
});
```

m.lock() is called

m.unlock() is called

Mutex RAII Wrappers

- ❖ Shared mutexes additionally need an RAII wrapper that calls `lock_shared` and `unlock_shared`
 - For this **`std::shared_lock`** can be used
 - Note
 - **`std::shared_lock`** is only movable and not copyable (unlike `std::shared_ptr`)

Example

```
std::shared_mutex m;
```

```
int i = 0;
```

```
std::thread t([&] {  
    std::shared_lock l{m};  
    std::cout << i;
```

```
    ...  
});
```

m.lock_shared() is called

m.unlock_shared() is
called

Avoiding Deadlocks using C++

- ❖ Deadlocks can occur when using multiple mutexes
 - When two different threads each succeed to lock a subset of the mutexes and then try to lock the rest
 - Can be avoided by always locking mutexes in a consistent order

Example

```
std::mutex m1, m2, m3;
```

```
void threadA() {  
    std::unique_lock l1{m1}, l2{m2}, l3{m3};  
}
```

```
void threadB() {  
    std::unique_lock l3{m3}, l2{m2}, l1{m1};  
    // DANGER: order not consistent with threadA()  
}
```

Concurrent calls to threadA() and threadB() can lead to deadlocks. E.g., A could get the locks for m1 and m2 while B gets a lock for m3.

Avoiding Deadlocks using C++

- ❖ Sometimes, it is not possible to always guarantee a consistent order
 - The function **std::lock** takes any number of mutexes and locks them all by using a deadlock-avoiding algorithm
 - **std::scoped_lock** is an RAII wrapper for **std::lock**

Example

```
std::mutex m1, m2, m3;
```

```
void threadA() {  
    std::scoped_lock l{m1, m2, m3};  
}
```

```
void threadB() {  
    std::scoped_lock l{m3, m2, m1};  
}
```

Here, calling threadA() and threadB() concurrently will not lead to deadlocks. Note: This should only be used if there is no other way as it is generally very inefficient!

Condition Variables

- ❖ A condition variable is a synchronization primitive that allows multiple threads to wait until an (arbitrary) condition becomes true
 - A condition variable uses a mutex to synchronize threads
 - Threads can wait on or notify the condition variable
 - When a thread waits on the condition variable, it blocks until another thread notifies it
 - If a thread waited on the condition variable and is notified, it holds the mutex
 - A notified thread must check the condition explicitly because spurious wake-ups can occur

Condition Variables

- ❖ The standard library defines the class `std::condition_variable` in the header `<condition_variable>` which has the following member functions
 - `wait()`
 - Takes a reference to a `std::unique_lock` that must be locked by the caller as an argument, unlocks the mutex and waits for the condition variable
 - `notify_one()`
 - Notify a single waiting thread, mutex does not need to be held by the caller
 - `notify_all()`
 - Notify all waiting threads, mutex does not need to be held by the caller

Example

```
using namespace std;

mutex m;
condition_variable cv;
queue<int> taskQueue;

void pushWork(int task) {
    unique_lock l{m};
    taskQueue.push(task);

    cv.notify_one();
}

void workerThread() {
    unique_lock l{m};

    while(true) {
        if (!taskQueue.empty())
        {
            int task = taskQueue.front();
            taskQueue.pop();
            l.unlock();
            // do something
            l.lock();
        }
        cv.wait(l);
    }
}
```

Tasks are inserted into a queue and then worker threads are notified to do the task

Atomic Operations

- ❖ Threads can also be synchronized with **atomic operations** that are usually much more efficient than mutexes
 - Atomicity means that an operation is executed as one unit, no other operation on the same object can be executed at the same time
 - This is usually implemented by using special CPU instructions, no operating system is needed
 - **Note**
 - Only the atomicity of single operations are guaranteed to be atomic

Example

If we could guarantee that `atomic_add` is performed in an atomic way, then `a = 13` will be the result nevertheless the scheduling of `threadA` and `threadB`

```
int a = 10;
```

```
void threadA() { atomic_add(a, 1); }  
void threadB() { atomic_add(a, 2); }
```

Atomic Operations Library

- ❖ All classes and functions related to atomic operations can be found in the `<atomic>` header
- ❖ `std::atomic<T>` is a class that represents an atomic version of the type T
- ❖ This class has several member functions that implement atomic operations
 - `T load()`
 - Loads the value (allows concurrent writes)
 - `void store(T desired)`
 - Stores desired in the object

Atomic Operations Library

- `T exchange(T desired)`
 - Stores desired in the object and returns the old value
- `bool compare_exchange_weak(...)` and `bool compare_exchange_strong(...)`
 - Perform a compare-and-swap (CAS) operation

Atomic Operations Library

- ❖ If T is an integral type, the following operations also exist
 - T fetch_add(T arg)
 - Adds arg to the value and returns the old value
 - T fetch_sub(T arg)
 - Same for subtraction
 - T fetch_and(T arg)
 - Same for bitwise and
 - T fetch_or(T arg)
 - Same for bitwise or
 - T fetch_xor(T arg)
 - Same for bitwise xor

Atomic Operations: Modification Order

- ❖ All modifications of a single atomic object are totally ordered in the so-called modification order
 - The modification order is consistent between threads (i.e. all threads see the same order)
 - The modification order is only total for individual atomic objects

Example

```
std::atomic<int> i = 0;
```

```
void workerThread() {  
    i.fetch_add(1); // (A)  
    i.fetch_sub(1); // (B)  
}
```

```
void readerThread() {  
    int iLocal = i.load();  
    if(iLocal == 0 || iLocal == 1)  
        std::cout << "Memory Consistent!";  
}
```

Because the memory order is consistent between threads, the reader thread will never see an execution order of (A), (B), (B), (A), for example.

This is always true